

HW5-Coding

January 7, 2024

1 Homework 5: Convolutional neural network (30 points)

In this part, you need to implement and train a convolutional neural network on the CIFAR-10 dataset with PyTorch. ### What is PyTorch?

PyTorch is a system for executing dynamic computational graphs over Tensor objects that behave similarly as numpy ndarray. It comes with a powerful automatic differentiation engine that removes the need for manual back-propagation.

1.0.1 Why?

- Our code will now run on GPUs! Much faster training. When using a framework like PyTorch or TensorFlow you can harness the power of the GPU for your own custom neural network architectures without having to write CUDA code directly (which is beyond the scope of this class).
- We want you to be ready to use one of these frameworks for your project so you can experiment more efficiently than if you were writing every feature you want to use by hand.
- We want you to stand on the shoulders of giants! TensorFlow and PyTorch are both excellent frameworks that will make your lives a lot easier, and now that you understand their guts, you are free to use them :)
- We want you to be exposed to the sort of deep learning code you might run into in academia or industry. ## How can I learn PyTorch?

Justin Johnson has made an excellent [tutorial](#) for PyTorch.

You can also find the detailed [API doc](#) here. If you have other questions that are not addressed by the API docs, the [PyTorch forum](#) is a much better place to ask than StackOverflow.

Install PyTorch and Skorch.

```
[1]: pip install -q torch skorch torchvision torchtext
```

```
[6]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import skorch
import sklearn
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

1.1 0. Tensor Operations (5 points)

Tensor operations are important in deep learning models. In this part, you are required to get familiar to some common tensor operations in PyTorch.

1.1.1 1) Tensor squeezing, unsqueezing and viewing

Tensor squeezing, unsqueezing and viewing are important methods to change the dimension of a Tensor, and the corresponding functions are [torch.squeeze](#), [torch.unsqueeze](#) and [torch.Tensor.view](#). Please read the documents of the functions, and finish the following practice.

```
[3]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2],
                  [3, 4],
                  [5, 6]])

x.shape
# Add two new dimensions to x by using the function torch.unsqueeze, so that
# the size of x becomes (3, 1, 2, 1).
x = torch.unsqueeze(torch.unsqueeze(x,1),-1)
print(x.shape)
# Remove the two dimensions just added by using the function torch.squeeze,
# and change the size of x back to (3, 2).
x = torch.squeeze(x)
print(x.shape)
# x is now a two-dimensional tensor, or in other words a matrix. Now use the
# function torch.Tensor.view and change x to a one-dimensional vector with
# size being (6).
x = x.view(6)
print(x.shape)
```

```
torch.Size([3, 1, 2, 1])
torch.Size([3, 2])
torch.Size([6])
```

1.1.2 2) Tensor concatenation and stack

Tensor concatenation and stack are operations to combine small tensors into big tensors. The corresponding functions are [torch.cat](#) and [torch.stack](#). Please read the documents of the functions, and finish the following practice.

```
[4]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2], [3, 4], [5, 6]])

# y is a tensor with size being (3, 2)
```

```

y = torch.Tensor([[[-1, -2], [-3, -4], [-5, -6]]])

# Our goal is to generate a tensor z with size as (2, 3, 2), and z[0,:,:] = x,
↳ z[1,:,:] = y.

# Use torch.stack to generate such a z
z = torch.stack((x,y))
print(z[0,:,:])
# Use torch.cat and torch.unsqueeze to generate such a z
z = torch.cat((torch.unsqueeze(x,0),torch.unsqueeze(y,0)))
print(z[1,:,:])

```

```

tensor([[1., 2.],
        [3., 4.],
        [5., 6.]])
tensor([[[-1., -2.],
        [-3., -4.],
        [-5., -6.]]])

```

1.1.3 3) Tensor expansion

Tensor expansion is to expand a tensor into a larger tensor along singleton dimensions. The corresponding functions are `torch.Tensor.expand` and `torch.Tensor.expand_as`. Please read the documents of the functions, and finish the following practice.

```

[5]: # x is a tensor with size being (3)
x = torch.Tensor([1, 2, 3])

# Our goal is to generate a tensor z with size (2, 3), so that z[0,:,:] = x,
↳ z[1,:,:] = x.

# [TO DO]
# Change the size of x into (1, 3) by using torch.unsqueeze.
x = torch.unsqueeze(x,0)
print(x.shape)

# [TO DO]
# Then expand the new tensor to the target tensor by using torch.Tensor.expand.
z = x.expand(2,3)
print(z.shape)

```

```

torch.Size([1, 3])
torch.Size([2, 3])

```

1.1.4 4) Tensor reduction in a given dimension

In deep learning, we often need to compute the mean/sum/max/min value in a given dimension of a tensor. Please read the document of `torch.mean`, `torch.sum`, `torch.max`, `torch.min`, `torch.topk`,

and finish the following practice.

```
[6]: # x is a random tensor with size being (10, 50)
x = torch.randn(10, 50)

# Compute the mean value for each row of x.
# You need to generate a tensor x_mean of size (10), and x_mean[k, :] is the
    ↪ mean value of the k-th row of x.
x_mean = torch.mean(x,1)
print(x_mean[3, ])

# Compute the sum value for each row of x.
# You need to generate a tensor x_sum of size (10).
x_sum = torch.sum(x,1)
print(x_sum.shape)

# Compute the max value for each row of x.
# You need to generate a tensor x_max of size (10).
x_max, s = torch.max(x,1)
print(x_max.shape)

# Compute the min value for each row of x.
# You need to generate a tensor x_min of size (10).
x_min, s = torch.min(x,1)
print(x_min.shape)

# Compute the top-5 values for each row of x.
# You need to generate a tensor x_mean of size (10, 5), and x_top[k, :] is the
    ↪ top-5 values of each row in x.
x_mean_xtop, s = torch.topk(x,5,1)
print((x_mean_xtop.shape))
```

```
tensor(0.1522)
torch.Size([10])
torch.Size([10])
torch.Size([10])
torch.Size([10, 5])
```

1.2 Convolutional Neural Networks

Implement a convolutional neural network for image classification on CIFAR-10 dataset.

CIFAR-10 is an image dataset of 10 categories. Each image has a size of 32x32 pixels. The following code will download the dataset, and split it into `train` and `test`. For this question, we use the default validation split generated by Skorch.

```
[3]: train = torchvision.datasets.CIFAR10("./data", train=True, download=True)
test = torchvision.datasets.CIFAR10("./data", train=False, download=True)
```

```
train.targets = torch.LongTensor(train.targets)
```

Files already downloaded and verified

Files already downloaded and verified

The following code visualizes some samples in the dataset. You may use it to debug your model if necessary.

```
[7]: def plot(data, labels=None, num_sample=5):
    n = min(len(data), num_sample)
    for i in range(n):
        plt.subplot(1, n, i+1)
        plt.imshow(data[i], cmap="gray")
        plt.xticks([])
        plt.yticks([])
        if labels is not None:
            plt.title(labels[i])

train.labels = [train.classes[target] for target in train.targets]
plot(train.data, train.labels)
```



1.2.1 1) Basic CNN implementation

Consider a basic CNN model

- It has 3 convolutional layers, followed by a linear layer.
- Each convolutional layer has a kernel size of 3, a padding of 1.
- ReLU activation is applied on every hidden layer.

Please implement this model in the following section. The hyperparameters is then be tuned and you need to fill the results in the table.

a) Implement convolutional layers (10 Points) Implement the initialization function and the forward function of the CNN.

```
[8]: class CNN(nn.Module):
    def __init__(self, channels):
        super(CNN, self).__init__()
        # implement parameter definitions here
```

```

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
self.conv1 = nn.Conv2d(3,channels,3,padding=1)
self.conv2 = nn.Conv2d(channels,channels,3,padding=1)
self.conv3 = nn.Conv2d(channels,channels,3,padding=1)
self.fcl = nn.Linear(channels*32*32,10)
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
def forward(self, images):
    # implement the forward function here
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    images = F.relu(self.conv1(images))
    images = F.relu(self.conv2(images))
    images = F.relu(self.conv3(images))
    images = images.view(images.size(0),-1)
    images = self.fcl(images)
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return images

```

b) **Tune hyperparameters** Train the CNN model on CIFAR-10 dataset. We can tune the number of channels, optimizer, learning rate and the number of epochs for best validation accuracy.

```

[10]: # implement hyperparameters, you can select and modify the hyperparameters by
      ↪yourself here.

optimize = [torch.optim.SGD, torch.optim.Adam]
learning_rate = [1e-5, 1e-4, 1e-3, 1e-2]
channel = [16,32,64]

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for o in optimize:
        for c in channel:
            print(f'The channel was {c}, the learning rate was {l} and the optimizer_
            ↪was {str(o)}')

            cnn = CNN(channels = c)

            model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
            ↪CrossEntropyLoss,

                                                device="cpu",
                                                optimizer=o,
                                                # optimizer__momentum=0.90,
                                                lr=l,
                                                max_epochs=25,

```

```

        batch_size=64,
        callbacks=[skorch.callbacks.
↳EarlyStopping(lower_is_better=True))]
        # implement input normalization & type cast here
        model.fit(train_data_normalized, np.asarray(train.targets))

```

The channel was 16, the learning rate was 1e-05 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.3031	0.0922	2.3030	
33.2964				
2	2.3030	0.0923	2.3030	
32.4063				
3	2.3029	0.0928	2.3029	
31.2097				
4	2.3029	0.0930	2.3028	
32.2927				
5	2.3028	0.0932	2.3028	
31.1628				
6	2.3027	0.0935	2.3027	
31.4149				
7	2.3027	0.0940	2.3026	
30.7325				
8	2.3026	0.0938	2.3026	30.8253
9	2.3025	0.0942	2.3025	
32.8183				
10	2.3025	0.0941	2.3024	32.3721
11	2.3024	0.0945	2.3024	
31.9081				
12	2.3024	0.0952	2.3023	
31.5482				
13	2.3023	0.0959	2.3022	
31.3438				
14	2.3022	0.0967	2.3022	
32.1400				
15	2.3022	0.0982	2.3021	
31.0039				
16	2.3021	0.0985	2.3020	
31.0092				
17	2.3020	0.0996	2.3020	
32.0398				
18	2.3020	0.1011	2.3019	
32.9767				
19	2.3019	0.1020	2.3018	
31.9393				
20	2.3018	0.1039	2.3018	

```

31.2194
  21      2.3018      0.1051      2.3017
32.5872
  22      2.3017      0.1059      2.3016
31.4957
  23      2.3016      0.1072      2.3016
31.6174
  24      2.3016      0.1090      2.3015
32.3750
  25      2.3015      0.1118      2.3014
32.3782
The channel was 32, the learning rate was 1e-05 and the optimizer was <class
'torch.optim.sgd.SGD'>

```

epoch	train_loss	valid_acc	valid_loss	dur
1	2.3026	0.1118	2.3025	
75.5154				
2	2.3023	0.1192	2.3022	
73.5765				
3	2.3021	0.1307	2.3019	
73.9044				
4	2.3018	0.1455	2.3016	
74.4299				
5	2.3015	0.1546	2.3014	
74.1025				
6	2.3013	0.1537	2.3011	74.7158
7	2.3010	0.1471	2.3008	75.1688
8	2.3008	0.1398	2.3006	74.0391
9	2.3005	0.1341	2.3004	74.8319
10	2.3003	0.1254	2.3001	78.7981
11	2.3001	0.1191	2.2999	74.6853
12	2.2998	0.1144	2.2996	75.1045
13	2.2996	0.1083	2.2994	75.2092
14	2.2994	0.1044	2.2992	75.3185
15	2.2992	0.1027	2.2990	73.9641
16	2.2990	0.1009	2.2988	75.9277
17	2.2988	0.0994	2.2985	74.4793
18	2.2985	0.0991	2.2983	81.8071
19	2.2983	0.0994	2.2981	74.3654
20	2.2981	0.1005	2.2979	75.5451
21	2.2979	0.1000	2.2977	74.9100
22	2.2977	0.1022	2.2975	75.4648
23	2.2975	0.1026	2.2973	74.9593
24	2.2973	0.1048	2.2971	75.1578
25	2.2971	0.1064	2.2968	74.4299

```

The channel was 64, the learning rate was 1e-05 and the optimizer was <class
'torch.optim.sgd.SGD'>

```

epoch	train_loss	valid_acc	valid_loss	dur
-------	------------	-----------	------------	-----

1	2.3036	0.1017	2.3032	
198.9562				
2	2.3030	0.1063	2.3026	
198.6153				
3	2.3024	0.1189	2.3021	
216.3862				
4	2.3019	0.1192	2.3015	
223.2036				
5	2.3014	0.1128	2.3010	216.7390
6	2.3009	0.1063	2.3005	217.7274
7	2.3005	0.1027	2.3001	217.4551
8	2.3000	0.1016	2.2996	204.9226
9	2.2996	0.1003	2.2992	188.9512
10	2.2991	0.1001	2.2987	187.3891
11	2.2987	0.1002	2.2983	190.4109
12	2.2982	0.0998	2.2978	189.4197
13	2.2978	0.0997	2.2974	188.3463
14	2.2974	0.0996	2.2969	187.0955
15	2.2969	0.0995	2.2965	189.2065
16	2.2965	0.0994	2.2960	191.9360
17	2.2961	0.0997	2.2956	189.1270
18	2.2956	0.0997	2.2951	191.8527
19	2.2952	0.0998	2.2946	194.3153
20	2.2947	0.1001	2.2941	195.2716
21	2.2942	0.1001	2.2937	195.4924
22	2.2937	0.1000	2.2932	193.3946
23	2.2933	0.1009	2.2927	196.2878
24	2.2928	0.1013	2.2922	195.4867
25	2.2923	0.1014	2.2916	190.2633

The channel was 16, the learning rate was 1e-05 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2091	0.2735	2.0809	
40.5973				
2	2.0070	0.3232	1.9368	
29.1095				
3	1.9032	0.3541	1.8627	
29.8156				
4	1.8488	0.3669	1.8235	
30.3466				
5	1.8154	0.3802	1.7928	
30.2068				
6	1.7868	0.3920	1.7661	
41.0151				
7	1.7605	0.4017	1.7412	
31.3739				

8	1.7365	0.4086	1.7184
31.1535			
9	1.7132	0.4165	1.6963
31.2500			
10	1.6912	0.4213	1.6774
42.8885			
11	1.6731	0.4237	1.6618
31.5856			
12	1.6574	0.4282	1.6478
31.1994			
13	1.6430	0.4318	1.6349
30.9401			
14	1.6297	0.4361	1.6229
32.1891			
15	1.6174	0.4398	1.6118
31.5050			
16	1.6059	0.4441	1.6014
30.8917			
17	1.5951	0.4468	1.5917
31.0846			
18	1.5849	0.4492	1.5825
77.5715			
19	1.5752	0.4510	1.5737
30.6983			
20	1.5659	0.4541	1.5653
31.0534			
21	1.5569	0.4564	1.5572
31.2100			
22	1.5483	0.4584	1.5494
31.0877			
23	1.5399	0.4624	1.5419
31.1569			
24	1.5318	0.4653	1.5347
31.5902			
25	1.5240	0.4679	1.5277
31.0664			

The channel was 32, the learning rate was 1e-05 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.1149	0.3102	1.9411	
74.5811				
2	1.8789	0.3659	1.8154	
75.3951				
3	1.7784	0.3966	1.7454	
75.3582				
4	1.7204	0.4144	1.6974	
76.2177				

5	1.6808	0.4245	1.6626
75.5505			
6	1.6494	0.4339	1.6335
74.8786			
7	1.6202	0.4427	1.6066
75.0219			
8	1.5946	0.4501	1.5847
75.0177			
9	1.5721	0.4590	1.5651
74.8159			
10	1.5513	0.4648	1.5471
75.3309			
11	1.5320	0.4698	1.5309
106.3320			
12	1.5143	0.4753	1.5164
101.5381			
13	1.4982	0.4784	1.5034
110.2510			
14	1.4836	0.4833	1.4916
74.4855			
15	1.4703	0.4884	1.4810
74.0732			
16	1.4580	0.4913	1.4713
75.0069			
17	1.4466	0.4930	1.4623
74.2623			
18	1.4360	0.4959	1.4540
74.7864			
19	1.4259	0.4992	1.4461
74.6959			
20	1.4163	0.5011	1.4385
72.3184			
21	1.4070	0.5018	1.4311
69.3383			
22	1.3980	0.5041	1.4240
71.3612			
23	1.3893	0.5057	1.4170
75.2446			
24	1.3807	0.5084	1.4101
75.8924			
25	1.3722	0.5103	1.4032
75.8757			

The channel was 64, the learning rate was 1e-05 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.0151	0.3576	1.8291	
202.1270				

2	1.7419	0.4183	1.6770
198.7369			
3	1.6265	0.4446	1.5937
201.3592			
4	1.5591	0.4619	1.5433
202.4845			
5	1.5155	0.4742	1.5105
201.6903			
6	1.4847	0.4827	1.4870
203.2900			
7	1.4607	0.4893	1.4681
203.0224			
8	1.4403	0.4940	1.4518
200.8325			
9	1.4219	0.4975	1.4366
202.5638			
10	1.4045	0.5019	1.4221
203.0319			
11	1.3877	0.5031	1.4080
201.4224			
12	1.3717	0.5075	1.3944
203.0768			
13	1.3564	0.5113	1.3812
201.1022			
14	1.3416	0.5164	1.3682
201.2386			
15	1.3271	0.5193	1.3552
202.8182			
16	1.3126	0.5223	1.3426
202.0798			
17	1.2989	0.5263	1.3308
203.6543			
18	1.2858	0.5293	1.3196
204.5527			
19	1.2733	0.5334	1.3090
202.7144			
20	1.2615	0.5381	1.2991
202.4044			
21	1.2502	0.5427	1.2899
202.2083			
22	1.2396	0.5456	1.2816
204.1165			
23	1.2296	0.5494	1.2738
204.8423			
24	1.2201	0.5541	1.2667
204.6857			
25	1.2112	0.5566	1.2601
198.3856			

The channel was 16, the learning rate was 0.0001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur

1	2.3032	0.1000	2.3029	
28.7028				
2	2.3028	0.1000	2.3026	28.3354
3	2.3025	0.1000	2.3024	28.4123
4	2.3023	0.1000	2.3021	27.5514
5	2.3021	0.0996	2.3019	28.1030
6	2.3019	0.1098	2.3017	
29.9903				
7	2.3017	0.1358	2.3015	
28.3431				
8	2.3014	0.1344	2.3012	28.8486
9	2.3012	0.1194	2.3010	28.3444
10	2.3009	0.1104	2.3007	29.0628
11	2.3006	0.1101	2.3004	29.2648
12	2.3003	0.1119	2.3000	28.6420
13	2.2999	0.1138	2.2996	29.5950
14	2.2994	0.1148	2.2991	29.9916
15	2.2990	0.1152	2.2985	31.6355
16	2.2984	0.1147	2.2979	29.3639
17	2.2978	0.1155	2.2973	29.2405
18	2.2971	0.1178	2.2965	29.2563
19	2.2963	0.1187	2.2957	28.8609
20	2.2955	0.1214	2.2948	28.9748
21	2.2945	0.1237	2.2937	29.7247
22	2.2934	0.1259	2.2925	29.3646
23	2.2921	0.1292	2.2911	29.2549
24	2.2907	0.1312	2.2895	28.1076
25	2.2890	0.1330	2.2876	29.5127

The channel was 32, the learning rate was 0.0001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur

1	2.3019	0.1010	2.3004	
71.7964				
2	2.2990	0.0957	2.2975	71.3975
3	2.2961	0.1035	2.2944	
72.1388				
4	2.2929	0.1090	2.2909	
71.7891				
5	2.2892	0.1129	2.2868	
72.7883				
6	2.2847	0.1206	2.2818	
71.2684				
7	2.2794	0.1338	2.2758	

70.9856			
8	2.2729	0.1507	2.2685
72.8750			
9	2.2650	0.1693	2.2595
71.1035			
10	2.2553	0.1898	2.2485
73.5009			
11	2.2432	0.2095	2.2347
72.6865			
12	2.2283	0.2266	2.2178
71.6683			
13	2.2100	0.2424	2.1972
71.6845			
14	2.1879	0.2544	2.1726
72.0716			
15	2.1620	0.2632	2.1445
71.5759			
16	2.1333	0.2672	2.1144
72.5669			
17	2.1039	0.2737	2.0849
71.9720			
18	2.0761	0.2800	2.0577
71.5721			
19	2.0511	0.2833	2.0338
71.3198			
20	2.0293	0.2873	2.0132
72.4818			
21	2.0106	0.2932	1.9956
73.3503			
22	1.9946	0.2983	1.9805
71.5927			
23	1.9807	0.3028	1.9673
71.5969			
24	1.9683	0.3082	1.9554
71.0115			
25	1.9570	0.3136	1.9445
71.7138			

The channel was 64, the learning rate was 0.0001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.3009	0.1120	2.2982	
195.1902				
2	2.2964	0.1066	2.2936	194.0137
3	2.2916	0.1109	2.2885	194.2230
4	2.2862	0.1200	2.2823	
195.4474				
5	2.2794	0.1364	2.2745	

195.2063			
6	2.2707	0.1554	2.2645
193.1715			
7	2.2595	0.1787	2.2514
194.5282			
8	2.2449	0.1982	2.2345
195.8728			
9	2.2262	0.2093	2.2133
194.8654			
10	2.2034	0.2224	2.1881
194.6705			
11	2.1772	0.2323	2.1603
195.2831			
12	2.1495	0.2421	2.1323
196.9269			
13	2.1227	0.2524	2.1065
193.8297			
14	2.0984	0.2648	2.0832
194.5961			
15	2.0761	0.2734	2.0616
194.1725			
16	2.0549	0.2848	2.0406
194.5781			
17	2.0340	0.2949	2.0197
196.1359			
18	2.0132	0.3010	1.9989
197.9209			
19	1.9929	0.3081	1.9790
195.3134			
20	1.9740	0.3152	1.9606
195.0304			
21	1.9568	0.3224	1.9441
197.6767			
22	1.9417	0.3277	1.9298
198.4896			
23	1.9287	0.3320	1.9175
196.9817			
24	1.9176	0.3345	1.9069
196.7081			
25	1.9080	0.3358	1.8977
193.1102			

The channel was 16, the learning rate was 0.0001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.9091	0.4104	1.6901	
2	1.6127	0.4607	1.5194	

30.9500			
3	1.5025	0.4793	1.4623
30.9233			
4	1.4517	0.4902	1.4305
30.2721			
5	1.4129	0.5029	1.3984
31.8046			
6	1.3701	0.5123	1.3735
31.1818			
7	1.3288	0.5216	1.3455
31.5237			
8	1.2937	0.5321	1.3172
31.0971			
9	1.2590	0.5437	1.2906
30.9779			
10	1.2283	0.5503	1.2708
30.2025			
11	1.2016	0.5568	1.2541
31.1930			
12	1.1771	0.5621	1.2402
31.1481			
13	1.1543	0.5667	1.2282
31.0480			
14	1.1332	0.5718	1.2180
30.6424			
15	1.1136	0.5761	1.2095
31.3306			
16	1.0954	0.5794	1.2020
31.3529			
17	1.0785	0.5811	1.1952
31.1834			
18	1.0627	0.5817	1.1891
31.1058			
19	1.0478	0.5851	1.1836
31.0362			
20	1.0336	0.5861	1.1784
30.9513			
21	1.0201	0.5883	1.1740
31.1171			
22	1.0072	0.5902	1.1700
31.0392			
23	0.9948	0.5936	1.1660
30.5617			
24	0.9829	0.5959	1.1625
30.8305			
25	0.9713	0.5972	1.1596
31.4569			

The channel was 32, the learning rate was 0.0001 and the optimizer was <class

'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7448	0.4688	1.5068	
75.5064				
2	1.4580	0.5062	1.3951	
75.1589				
3	1.3332	0.5375	1.3080	
75.4120				
4	1.2465	0.5582	1.2512	
75.9160				
5	1.1877	0.5734	1.2123	
75.0868				
6	1.1410	0.5847	1.1852	
74.9629				
7	1.1002	0.5949	1.1653	
75.5087				
8	1.0633	0.5996	1.1492	
75.9712				
9	1.0295	0.6058	1.1342	
76.0611				
10	0.9980	0.6107	1.1192	
75.7848				
11	0.9679	0.6156	1.1054	
75.7735				
12	0.9389	0.6199	1.0942	
76.3694				
13	0.9105	0.6219	1.0854	
76.6931				
14	0.8830	0.6250	1.0779	
75.8649				
15	0.8561	0.6296	1.0724	
76.3687				
16	0.8300	0.6347	1.0692	
76.4422				
17	0.8044	0.6380	1.0671	
83.2052				
18	0.7794	0.6387	1.0675	75.9804
19	0.7552	0.6392	1.0690	75.9325
20	0.7316	0.6395	1.0721	76.2727
21	0.7087	0.6395	1.0767	76.2571

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 64, the learning rate was 0.0001 and the optimizer was <class

'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6846	0.4873	1.4493	
199.5793				

2	1.3539	0.5523	1.2738	
199.9522				
3	1.2091	0.5765	1.2135	
200.8681				
4	1.1118	0.5935	1.1575	
200.6454				
5	1.0238	0.6135	1.1113	
201.8104				
6	0.9444	0.6264	1.0786	
201.4504				
7	0.8754	0.6313	1.0595	
203.6201				
8	0.8145	0.6402	1.0516	
202.8307				
9	0.7593	0.6422	1.0526	200.6320
10	0.7077	0.6450	1.0609	203.2361
11	0.6587	0.6443	1.0765	203.4782
12	0.6113	0.6437	1.0997	352.7898

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 16, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2964	0.1563	2.2855	
62.8376				
2	2.2500	0.2369	2.1857	
63.1041				
3	2.0977	0.2771	2.0253	
61.0745				
4	1.9956	0.3003	1.9668	
58.8770				
5	1.9464	0.3193	1.9265	
59.4622				
6	1.9105	0.3325	1.8931	
61.0843				
7	1.8791	0.3464	1.8611	
59.7554				
8	1.8484	0.3598	1.8307	
59.1445				
9	1.8205	0.3675	1.8045	
59.5986				
10	1.7975	0.3725	1.7836	
55.7545				
11	1.7786	0.3777	1.7670	
52.4581				
12	1.7627	0.3846	1.7533	
57.9407				
13	1.7489	0.3882	1.7418	

59.9380			
14	1.7367	0.3925	1.7317
65.5651			
15	1.7257	0.3944	1.7224
57.1700			
16	1.7155	0.3960	1.7138
31.3681			
17	1.7055	0.3986	1.7051
35.0293			
18	1.6956	0.4018	1.6961
39.7325			
19	1.6852	0.4064	1.6861
38.0291			
20	1.6740	0.4115	1.6748
36.4835			
21	1.6617	0.4163	1.6620
41.1977			
22	1.6481	0.4227	1.6480
41.3113			
23	1.6336	0.4269	1.6332
34.2660			
24	1.6189	0.4337	1.6187
39.6896			
25	1.6043	0.4382	1.6046
39.7966			

The channel was 32, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2918	0.1775	2.2719	
100.1234				
2	2.1953	0.2770	2.0722	
98.2262				
3	1.9944	0.3046	1.9449	
103.1981				
4	1.9145	0.3318	1.8852	
104.1640				
5	1.8644	0.3492	1.8420	
104.9741				
6	1.8289	0.3612	1.8109	
105.7524				
7	1.8018	0.3664	1.7870	
105.2751				
8	1.7796	0.3738	1.7674	
102.9188				
9	1.7605	0.3808	1.7508	
100.9611				
10	1.7439	0.3855	1.7362	

105.1222			
11	1.7285	0.3919	1.7221
105.9550			
12	1.7134	0.3981	1.7078
106.2921			
13	1.6978	0.4039	1.6923
105.0269			
14	1.6808	0.4117	1.6747
101.3276			
15	1.6623	0.4192	1.6553
106.9754			
16	1.6433	0.4245	1.6361
106.4313			
17	1.6254	0.4321	1.6191
106.3367			
18	1.6094	0.4386	1.6042
105.7679			
19	1.5948	0.4425	1.5906
105.5507			
20	1.5807	0.4490	1.5774
79.0545			
21	1.5669	0.4531	1.5643
110.3573			
22	1.5530	0.4584	1.5511
111.5549			
23	1.5390	0.4614	1.5379
112.9123			
24	1.5251	0.4657	1.5249
113.7903			
25	1.5112	0.4707	1.5122

111.5942

The channel was 64, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.2575	0.2347	2.1720	
284.8710				
2	2.0469	0.3039	1.9570	
271.1822				
3	1.9155	0.3392	1.8792	
257.0626				
4	1.8564	0.3573	1.8284	
252.7677				
5	1.8129	0.3730	1.7880	
253.8607				
6	1.7745	0.3881	1.7490	
303.3023				
7	1.7372	0.4042	1.7110	

222.1229			
8	1.7038	0.4136	1.6800
241.7212			
9	1.6768	0.4232	1.6558
202.0839			
10	1.6538	0.4310	1.6353
263.5292			
11	1.6329	0.4373	1.6168
265.5118			
12	1.6132	0.4438	1.5992
251.3128			
13	1.5942	0.4494	1.5822
226.1318			
14	1.5756	0.4549	1.5653
241.7881			
15	1.5574	0.4595	1.5487
192.4584			
16	1.5395	0.4634	1.5325
194.0290			
17	1.5223	0.4689	1.5173
198.1915			
18	1.5060	0.4732	1.5035
214.3901			
19	1.4909	0.4765	1.4911
214.4707			
20	1.4770	0.4790	1.4797
213.0608			
21	1.4642	0.4803	1.4692
212.3046			
22	1.4521	0.4821	1.4595
213.4897			
23	1.4405	0.4861	1.4501
213.9201			
24	1.4294	0.4909	1.4412
213.1770			
25	1.4184	0.4937	1.4325
214.2598			

The channel was 16, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7436	0.4735	1.4827	
33.6238				
2	1.4139	0.5314	1.3268	
34.9068				
3	1.2470	0.5739	1.2012	
34.2287				
4	1.1048	0.5939	1.1564	

```

33.9101
  5      0.9872      0.6057      1.1545
33.7875
  6      0.8806      0.6068      1.1838  33.9868
  7      0.7912      0.6066      1.2303  33.8943
  8      0.7121      0.6018      1.3180  33.9959
  9      0.6346      0.5843      1.5201  34.5033

```

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 32, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6392	0.5129	1.3716	
79.5216				
2	1.2768	0.5762	1.1864	
80.1609				
3	1.0843	0.5977	1.1543	
81.0517				
4	0.9439	0.6044	1.1950	80.9044
5	0.8090	0.6082	1.2552	81.2396
6	0.6837	0.5862	1.4491	80.4444
7	0.5770	0.5707	1.6861	80.3253

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 64, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.5455	0.5606	1.2410	
215.2249				
2	1.1250	0.6300	1.0641	
218.1638				
3	0.9070	0.6414	1.0592	
218.0134				
4	0.7451	0.6298	1.1461	220.1412
5	0.6064	0.6173	1.3049	217.9364
6	0.4885	0.6053	1.4677	218.7542
7	0.4168	0.6066	1.5865	220.4605

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 16, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.0554	0.3064	1.9374	
32.1266				
2	1.7919	0.3784	1.7518	
31.9039				
3	1.7046	0.4043	1.6787	
32.3028				

4	1.6251	0.4372	1.5879
31.8986			
5	1.5426	0.4630	1.5165
32.4555			
6	1.4749	0.4854	1.4505
32.0101			
7	1.4209	0.5007	1.3991
32.1931			
8	1.3768	0.5098	1.3647
32.3143			
9	1.3407	0.5189	1.3422
34.0823			
10	1.3103	0.5260	1.3247
32.7837			
11	1.2837	0.5310	1.3127
32.8594			
12	1.2589	0.5351	1.3050
32.8592			
13	1.2350	0.5375	1.2950
32.2958			
14	1.2113	0.5429	1.2863
32.6479			
15	1.1864	0.5448	1.2753
32.5674			
16	1.1600	0.5510	1.2620
32.5912			
17	1.1338	0.5579	1.2510
32.1261			
18	1.1076	0.5624	1.2430
33.1683			
19	1.0813	0.5708	1.2259
32.5658			
20	1.0546	0.5763	1.2112
32.9377			
21	1.0271	0.5802	1.2013
32.5382			
22	0.9991	0.5859	1.1914
32.8069			
23	0.9713	0.5882	1.1885
33.4357			
24	0.9440	0.5933	1.1867
32.8244			
25	0.9175	0.5937	1.1897
33.5969			

The channel was 32, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.9882	0.3503	1.8243	

78.8034				
2	1.7197	0.4227	1.6394	
79.1140				
3	1.5920	0.4604	1.5231	
78.4201				
4	1.4871	0.4865	1.4378	
77.9095				
5	1.4048	0.5134	1.3671	
79.3930				
6	1.3368	0.5274	1.3144	
78.3532				
7	1.2702	0.5384	1.2814	
78.9115				
8	1.2169	0.5473	1.2562	
78.8646				
9	1.1693	0.5588	1.2340	
79.3294				
10	1.1236	0.5699	1.2112	
79.0200				
11	1.0788	0.5806	1.1873	
79.5395				
12	1.0351	0.5904	1.1709	
79.4168				
13	0.9931	0.5929	1.1610	
80.7001				
14	0.9530	0.6008	1.1532	
79.4426				
15	0.9143	0.6039	1.1506	
79.4936				
16	0.8764	0.6059	1.1516	79.8378
17	0.8392	0.6081	1.1577	81.5592
18	0.8021	0.6086	1.1680	79.7244
19	0.7651	0.6106	1.1833	80.8649

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 64, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.9238	0.3923	1.7129	
213.7789				
2	1.6223	0.4537	1.5411	
216.7653				
3	1.4901	0.4887	1.4371	
216.1997				
4	1.4097	0.5048	1.3796	
218.8107				
5	1.3467	0.5199	1.3309	
216.5749				

6	1.2840	0.5423	1.2776	
214.9383				
7	1.2187	0.5589	1.2334	
215.9890				
8	1.1613	0.5735	1.1945	
216.7445				
9	1.1131	0.5865	1.1653	
216.2992				
10	1.0708	0.5958	1.1396	
217.9955				
11	1.0329	0.6027	1.1224	
228.4997				
12	0.9984	0.6096	1.1077	
224.0026				
13	0.9654	0.6143	1.0994	
216.9344				
14	0.9331	0.6200	1.0930	
236.4126				
15	0.9002	0.6204	1.0890	
247.5127				
16	0.8663	0.6213	1.0879	
242.1076				
17	0.8310	0.6250	1.0874	
226.0298				
18	0.7947	0.6259	1.0890	216.5611
19	0.7578	0.6282	1.0964	217.5583
20	0.7206	0.6294	1.1091	254.2584
21	0.6830	0.6274	1.1272	260.1418

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 16, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.3163	0.1000	2.3034	
50.7011				
2	2.3038	0.1000	2.3035	47.4422
3	2.3039	0.1000	2.3035	42.5227
4	2.3039	0.1000	2.3035	41.7330
5	2.3039	0.1000	2.3035	38.6801

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 32, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.3153	0.1000	2.3035	
89.8181				
2	2.3039	0.1000	2.3035	90.8093
3	2.3039	0.1000	2.3035	123.9652

4	2.3039	0.1000	2.3035	135.9307
5	2.3039	0.1000	2.3035	118.6586

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 64, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.7192	0.1000	2.3034	
262.2795				
2	2.3038	0.1000	2.3035	242.4639
3	2.3039	0.1000	2.3035	280.0173
4	2.3039	0.1000	2.3035	268.4130
5	2.3039	0.1000	2.3035	440.0961

Stopping since valid_loss has not improved in the last 5 epochs.

Write down **validation accuracy** of your model under different hyperparameter settings. Note the validation set is automatically split by Skorch during `model.fit()`.

#channel for each layer	optimizer	SGD	Adam
(16,16,16)		62.64	66.00
(32,32,32)		67.89	72.11
(64,64,64)		63.22	65.60

1.2.2 2) Full CNN implementation (10 points)

Based on the CNN in the previous question, implement a full CNN model with max pooling layer.

- Add a max pooling layer after each convolutional layer.
- Each max pooling layer has a kernel size of 2 and a stride of 2.

Please implement this model in the following section. The hyperparameters is then be tuned and fill the results in the table. You are also required to complete the questions.

a) Implement max pooling layers Similar to the CNN implementation in previous question, implement max pooling layers.

```
[9]: class CNN_MaxPool(nn.Module):
    def __init__(self, channels):
        super(CNN_MaxPool, self).__init__()
        # implement parameter definitions here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        self.conv1 = nn.Conv2d(3,512,3,padding=1)
        self.conv2 = nn.Conv2d(512,512,3,padding=1)
        self.conv3 = nn.Conv2d(512,512,3,padding=1)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fcl = nn.Linear(512*4*4,10)
        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    def forward(self, images):
```

```

# implement the forward function here
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
images = self.pool(F.relu(self.conv1(images)))
images = self.pool(F.relu(self.conv2(images)))
images = self.pool(F.relu(self.conv3(images)))
images = images.view(images.size(0), -1)
images = self.fcl(images)
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
return images

```

b) **Tune hyperparameters** Based on the better optimizer found in the previous problem, we can tune the number of channels and learning rate for best validation accuracy.

```

[7]: # implement hyperparameters, you can select and modify the hyperparameters by
      ↪yourself here.
learning_rate = [1e-5, 1e-4, 1e-3, 1e-2]
channel = [16,32,64]
# Select the better optimizer by the result shown in the previous problem, you
      ↪can select and modify it by yourself here.
better_optimizer = torch.optim.Adam

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for c in channel:
        print(f'The channel was {c}, the learning rate was {l}')

        cnn = CNN_MaxPool(channels = c)

        model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
      ↪CrossEntropyLoss,
                                           device="cpu",
                                           optimizer=better_optimizer,
                                           lr=l,
                                           max_epochs=25,
                                           batch_size=64,
                                           callbacks=[skorch.callbacks.
      ↪EarlyStopping(lower_is_better=True)])
        # implement input normalization & type cast here
        model.fit(train_data_normalized, np.asarray(train.targets))

```

The channel was 16, the learning rate was 1e-05

epoch	train_loss	valid_acc	valid_loss	dur
1	1.9620	0.4077	1.7026	
2015.3576				

2	1.6225	0.4595	1.5351
2181.7414			
3	1.4975	0.4886	1.4461
2189.5115			
4	1.4178	0.5118	1.3862
2183.9949			
5	1.3609	0.5283	1.3418
2181.5302			
6	1.3178	0.5435	1.3058
2182.7197			
7	1.2821	0.5561	1.2747
2183.1216			
8	1.2505	0.5653	1.2466
2175.7946			
9	1.2216	0.5737	1.2212
2122.4000			
10	1.1948	0.5821	1.1977
2116.9122			
11	1.1697	0.5895	1.1760
2111.7452			
12	1.1462	0.5974	1.1558
2133.8165			
13	1.1239	0.6046	1.1369
2165.6519			
14	1.1029	0.6119	1.1189
2179.4310			
15	1.0829	0.6195	1.1020
2176.8308			
16	1.0637	0.6238	1.0863
2186.6272			
17	1.0455	0.6263	1.0712
2190.6645			
18	1.0280	0.6311	1.0570
2132.4965			
19	1.0113	0.6364	1.0436
2491.9667			
20	0.9951	0.6418	1.0309
2022.1192			
21	0.9796	0.6457	1.0188
1885.7328			
22	0.9648	0.6514	1.0072
1745.3710			
23	0.9505	0.6555	0.9963
1745.1618			
24	0.9367	0.6584	0.9858
1770.0926			
25	0.9235	0.6600	0.9759
2212.5278			

The channel was 32, the learning rate was 1e-05				
epoch	train_loss	valid_acc	valid_loss	dur

1	1.9730	0.4036	1.7164	
1949.7775				
2	1.6381	0.4532	1.5527	
1489.9995				
3	1.5127	0.4858	1.4591	
1483.2545				
4	1.4310	0.5029	1.3969	
1962.5577				
5	1.3724	0.5189	1.3522	
1624.4438				
6	1.3286	0.5328	1.3165	
1860.1606				
7	1.2928	0.5456	1.2858	
2097.1347				
8	1.2615	0.5581	1.2582	
1860.9127				
9	1.2330	0.5702	1.2329	
1925.6908				
10	1.2065	0.5782	1.2097	
1926.7113				
11	1.1819	0.5853	1.1880	
1904.7657				
12	1.1586	0.5928	1.1679	
2011.2623				
13	1.1366	0.6000	1.1489	
1861.9987				
14	1.1157	0.6052	1.1312	
1962.1067				
15	1.0958	0.6104	1.1144	
2043.5940				
16	1.0768	0.6162	1.0988	
2634.7108				
17	1.0587	0.6199	1.0841	
2675.6508				
18	1.0413	0.6253	1.0700	
2109.1479				
19	1.0246	0.6296	1.0566	
2029.2728				
20	1.0086	0.6341	1.0440	
2033.8819				
21	0.9933	0.6392	1.0322	
2029.2635				
22	0.9785	0.6440	1.0208	
2028.9265				
23	0.9644	0.6473	1.0100	

2025.9630				
24	0.9507	0.6506	0.9996	
2036.2984				
25	0.9375	0.6535	0.9899	
2012.0991				
The channel was 64, the learning rate was 1e-05				
epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.9789	0.4005	1.7277	
1985.6530				
2	1.6450	0.4538	1.5547	
1985.4058				
3	1.5141	0.4847	1.4577	
1987.3407				
4	1.4299	0.5104	1.3948	
1992.3526				
5	1.3713	0.5266	1.3502	
1986.3687				
6	1.3274	0.5381	1.3144	
2005.1232				
7	1.2912	0.5513	1.2838	
2030.5202				
8	1.2592	0.5605	1.2562	
2034.0302				
9	1.2302	0.5709	1.2308	
2040.9535				
10	1.2033	0.5802	1.2074	
3005.5707				
11	1.1782	0.5881	1.1854	
3138.7492				
12	1.1547	0.5942	1.1648	
2517.8228				
13	1.1324	0.6005	1.1459	
1770.9409				
14	1.1113	0.6077	1.1280	
1777.9515				
15	1.0913	0.6124	1.1113	
1827.6349				
16	1.0723	0.6176	1.0952	
1772.7285				
17	1.0541	0.6233	1.0803	
1825.2955				
18	1.0368	0.6281	1.0660	
1796.1912				
19	1.0201	0.6321	1.0526	
1692.3421				
20	1.0042	0.6363	1.0400	
1692.9635				

21	0.9889	0.6421	1.0281	
1839.3732				
22	0.9742	0.6457	1.0168	
2211.6594				
23	0.9601	0.6494	1.0062	
1855.1671				
24	0.9465	0.6535	0.9959	
1824.7757				
25	0.9334	0.6573	0.9862	
1802.4270				
The channel was 16, the learning rate was 0.0001				
epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6355	0.5304	1.3180	
2770.2951				
2	1.2318	0.6103	1.1098	
2611.8159				
3	1.0593	0.6553	0.9850	
1812.2924				
4	0.9358	0.6780	0.9174	
1520.0275				
5	0.8430	0.6907	0.8806	
2006.2764				
6	0.7693	0.7014	0.8549	
1987.1612				
7	0.7071	0.7083	0.8357	
2720.4578				
8	0.6510	0.7160	0.8158	
2284.5833				
9	0.5984	0.7265	0.7918	
1764.6349				
10	0.5477	0.7343	0.7780	
1683.3413				
11	0.4995	0.7381	0.7768	
1904.5913				
12	0.4555	0.7387	0.7776	1866.6369
13	0.4127	0.7384	0.7880	1862.4910
14	0.3709	0.7325	0.8154	1875.5805
15	0.3301	0.7300	0.8527	2020.2487

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 32, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6505	0.5274	1.3286	
2016.2083				
2	1.2360	0.6101	1.1119	
2046.6934				
3	1.0577	0.6581	0.9798	

2024.6255				
4	0.9334	0.6789	0.9123	
2022.7825				
5	0.8398	0.6931	0.8746	
2029.5375				
6	0.7639	0.7020	0.8480	
2022.9922				
7	0.6994	0.7123	0.8289	
1999.5237				
8	0.6407	0.7202	0.8144	
1841.0888				
9	0.5853	0.7250	0.8062	
1830.9262				
10	0.5335	0.7293	0.8037	
1830.7831				
11	0.4839	0.7342	0.8082	1844.9272
12	0.4376	0.7346	0.8202	1890.5180
13	0.3938	0.7356	0.8305	1852.2313
14	0.3516	0.7374	0.8523	1855.6431

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 64, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6386	0.5373	1.3020	
2788.4585				
2	1.2156	0.6161	1.0923	
1982.6056				
3	1.0390	0.6598	0.9749	
1952.9520				
4	0.9179	0.6826	0.9064	
1991.6278				
5	0.8271	0.6940	0.8686	
1930.3556				
6	0.7540	0.7066	0.8407	
2052.6165				
7	0.6913	0.7151	0.8192	
1921.7194				
8	0.6339	0.7215	0.8040	
1912.6408				
9	0.5806	0.7259	0.7941	
1920.7272				
10	0.5306	0.7316	0.7892	
1884.8599				
11	0.4836	0.7348	0.7915	2832.6945
12	0.4400	0.7368	0.7914	2918.9338
13	0.3984	0.7416	0.7927	2502.4152
14	0.3583	0.7412	0.8081	2751.9188

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 16, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
1	1.5403	0.5760	1.1992	
2208.2255				
2	1.1425	0.6306	1.0490	
1908.7850				
3	0.9650	0.6673	0.9531	
2381.4564				
4	0.8508	0.6886	0.9051	
2243.9538				
5	0.7680	0.6850	0.9369	3039.4427
6	0.6954	0.6813	0.9643	2009.9776
7	0.6422	0.6800	0.9796	1852.9855
8	0.5855	0.6866	0.9889	1923.2216

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 32, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
1	1.5969	0.5568	1.2330	
1852.2973				
2	1.1545	0.6319	1.0502	
1919.2373				
3	0.9796	0.6580	0.9890	
1955.3820				
4	0.8679	0.6725	0.9481	
2065.3813				
5	0.7888	0.6797	0.9225	
2047.5048				
6	0.7265	0.6810	0.9570	2038.7800
7	0.6762	0.6782	0.9940	2038.7441
8	0.6333	0.6657	1.0982	2026.2662
9	0.5929	0.6675	1.0986	2027.5202

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 64, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
1	1.6344	0.5295	1.3157	
1943.4180				
2	1.2120	0.6068	1.1103	
1900.2877				
3	1.0708	0.6285	1.0560	
1920.2538				
4	0.9739	0.6414	1.0254	
2047.0268				
5	0.8986	0.6423	1.0417	2594.3312
6	0.8384	0.6337	1.0671	2910.5896
7	0.7984	0.6396	1.0678	2163.7094

8 0.7519 0.6355 1.1093 1508.6315
 Stopping since valid_loss has not improved in the last 5 epochs.
 The channel was 16, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
1	5.3360	0.1000	2.3033	1572.3764
2	2.3038	0.1000	2.3034	1639.0350
3	2.3038	0.1000	2.3035	1821.6766
4	2.3039	0.1000	2.3035	1756.7726
5	2.3039	0.1000	2.3035	1845.7966

Stopping since valid_loss has not improved in the last 5 epochs.
 The channel was 32, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
1	4.6576	0.4688	1.4754	1711.5144
2	1.4273	0.5073	1.3865	1621.4952
3	1.3586	0.5256	1.3293	1622.3163

Write down the **validation accuracy** of the model under different hyperparameter settings.

#channel for each layer	validation accuracy
(16,16,16)	75.20
(32,32,32)	75.13
(64,64,64)	79.72

For the best model you have, test it on the test set.

```
[ ]: # implement the same input normalization & type cast here
test_data_normalized = torch.Tensor(test.data/255)
test_data_normalized = test_data_normalized.permute(0,3,1,2)
test.predictions = model.predict(test_data_normalized)
sklearn.metrics.accuracy_score(test.targets, test.predictions)
```

How much **test accuracy** do you get? What can you conclude for the design of CNN structure and tuning of hyperparameters? (5 points)

Your Answer:* I get 70.5% test accuracy. The number of channels and learning rate are the most important hyperparameters. They should not be too large or too small. Adam optimizer is better than SGD optimizer because it can converge faster. The structure of CNN is also important. The max pooling layer can improve the performance of CNN. The kernel size and stride of max pooling layer should be appropriate.*