



Introduction to Machine Learning CS182

Lu Sun

School of Information Science and Technology

ShanghaiTech University

December 28, 2023

Today:

- Deep Generative Networks (DGN)

Readings:


- Deep Learning (DL), Chapters 14&20

Today's Agenda

Deep Generative Models (DGM)

- Overview
- Representation Learning with Autoencoder
- Generative Adversarial Network (GAN)
- Applications of GANs

Acknowledgement: Hugo Larochelle's, MehryarMohri@NYU's, YingyuLiang@Princeton's, BhikshaRaj@CMU's&FeifeiLi@Stanford's course notes



Deep Generative Networks-DGM

- **Overview**
 - Representation Learning with Autoencoder
 - Generative Adversarial Network (GAN)
 - Applications of GANs
- 

Unsupervised Learning

- Task formulation

Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.

Unsupervised Learning

- Task formulation

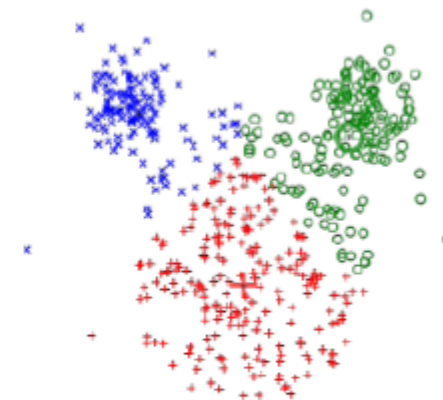
Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



K-means clustering

Unsupervised Learning

■ Task formulation

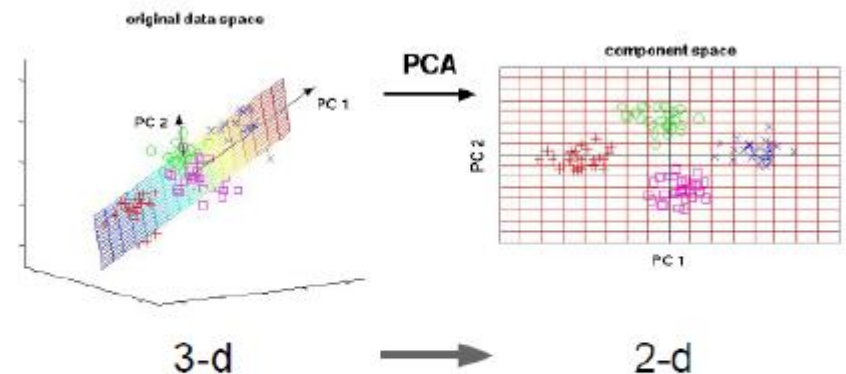
Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



Principal Component Analysis
(Dimensionality reduction)

Unsupervised Learning

■ Task formulation

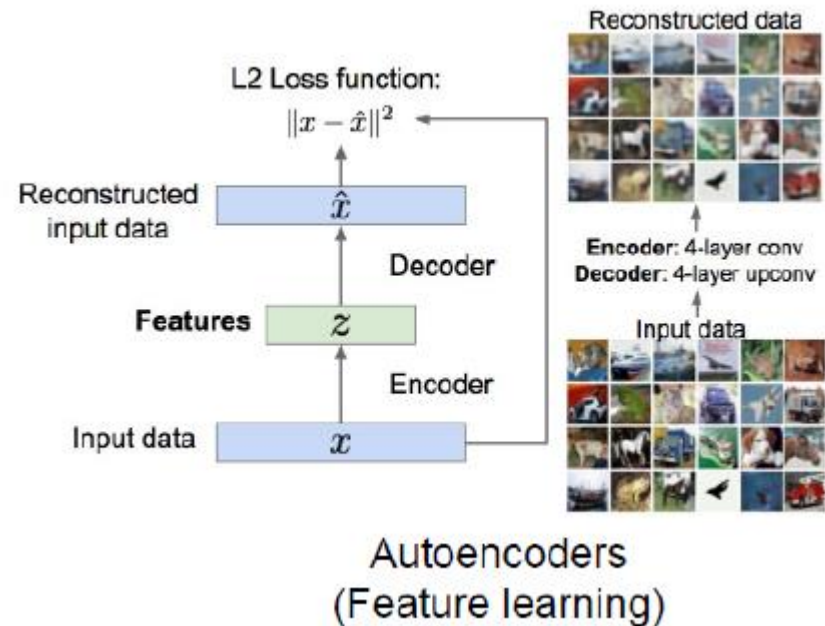
Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



Unsupervised Learning

■ Task formulation

Unsupervised Learning

Data: x

Just data, no labels!



Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.




Figure copyright anandadurai, 2016. Reproduced with permission.


1-d density estimation



2-d density estimation



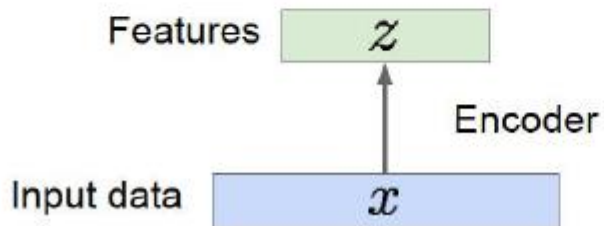
Deep Generative Networks-DGM

- Overview
 - **Representation Learning with Autoencoder**
 - Generative Adversarial Network (GAN)
 - Applications of GANs
- 

Autoencoder

■ Feature representation learning

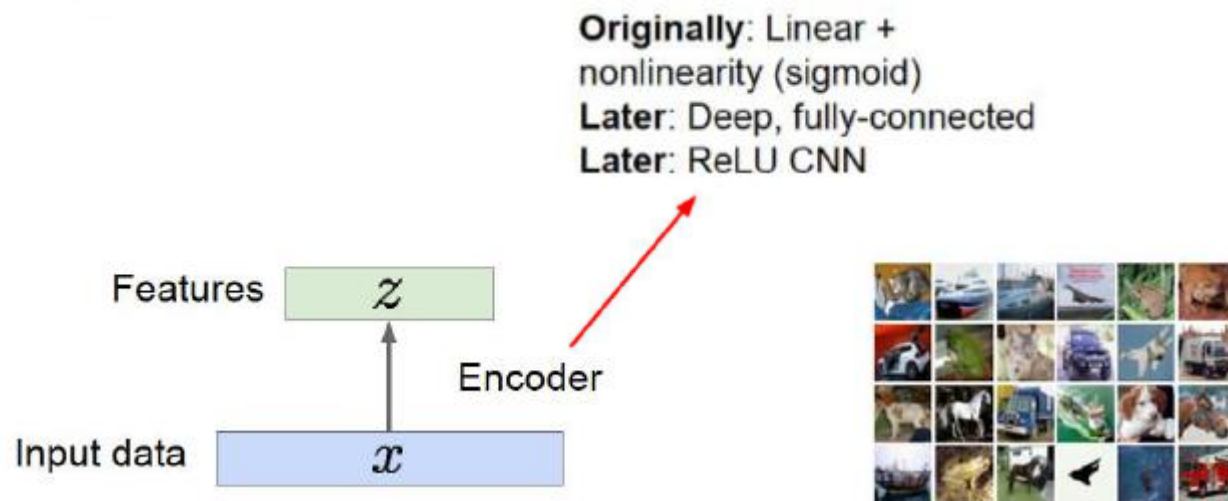
Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



Autoencoder

■ Feature representation learning

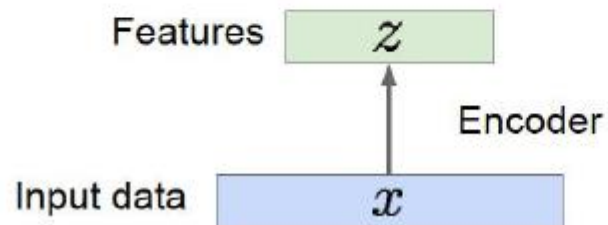
Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



Autoencoder

■ Feature representation learning

How to learn this feature representation?

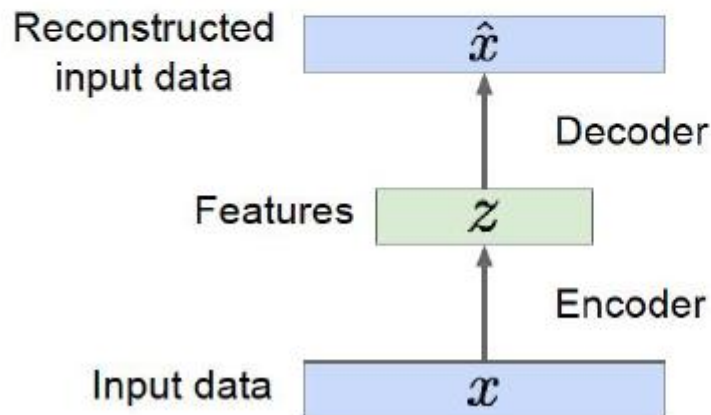


Autoencoder

■ Feature representation learning

How to learn this feature representation?

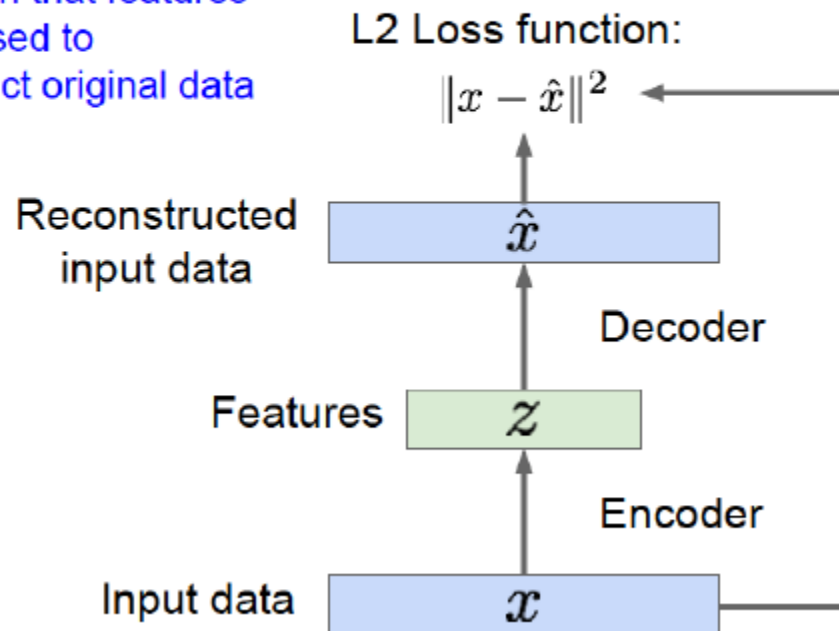
Train such that features can be used to reconstruct original data
"Autoencoding" - encoding itself



Autoencoder

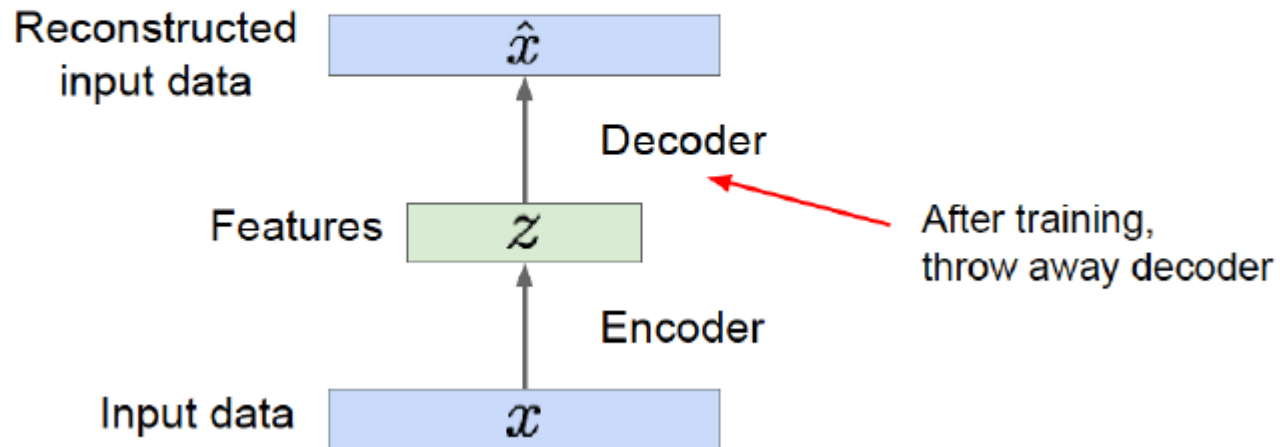
■ Feature representation learning

Train such that features
can be used to
reconstruct original data



Autoencoder

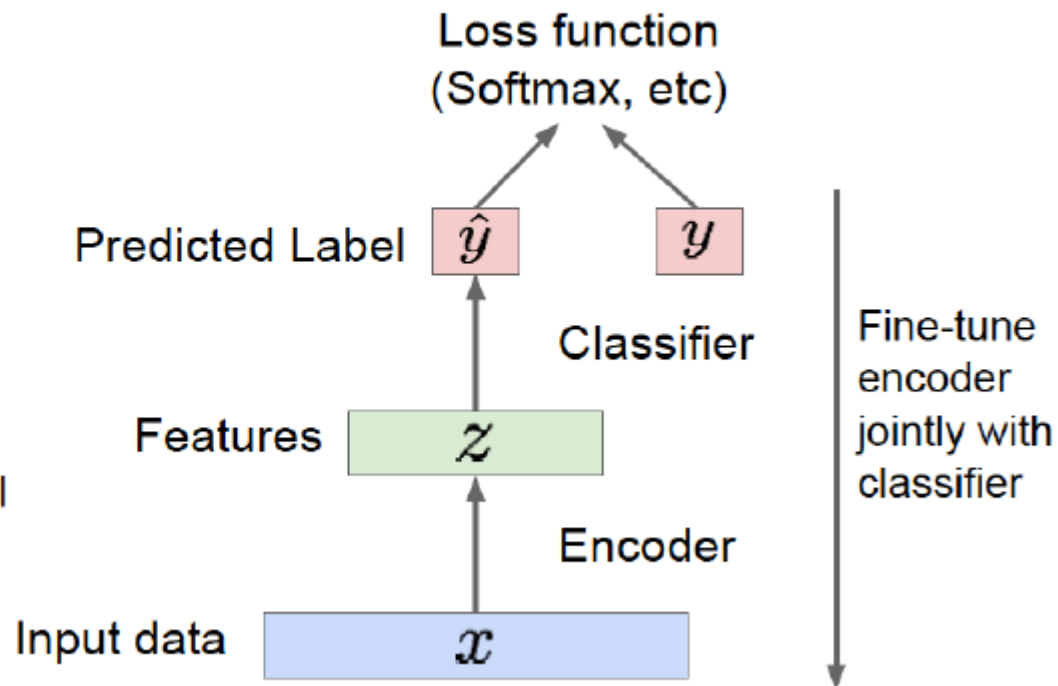
- Feature representation learning



Autoencoder

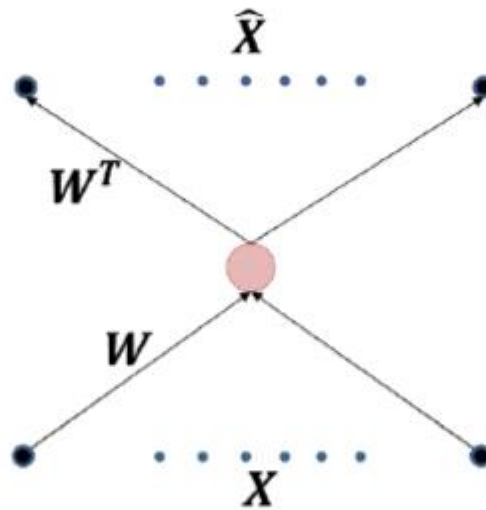
- Feature representation learning

Encoder can be used to initialize a **supervised** model



Autoencoder


■ Linear hidden layer example



- A single hidden unit
- Hidden unit has *linear* activation
- What will this learn?

Autoencoder

■ Linear hidden layer example



Training: Learning W by minimizing L2 divergence

$$\hat{x} = w^T w x$$

$$\text{div}(\hat{x}, x) = \|x - \hat{x}\|^2 = \|x - w^T w x\|^2$$


$$\hat{W} = \underset{W}{\operatorname{argmin}} E[\text{div}(\hat{x}, x)]$$

$$\hat{W} = \underset{W}{\operatorname{argmin}} E[\|x - w^T w x\|^2]$$

- This is just PCA!

Autoencoder

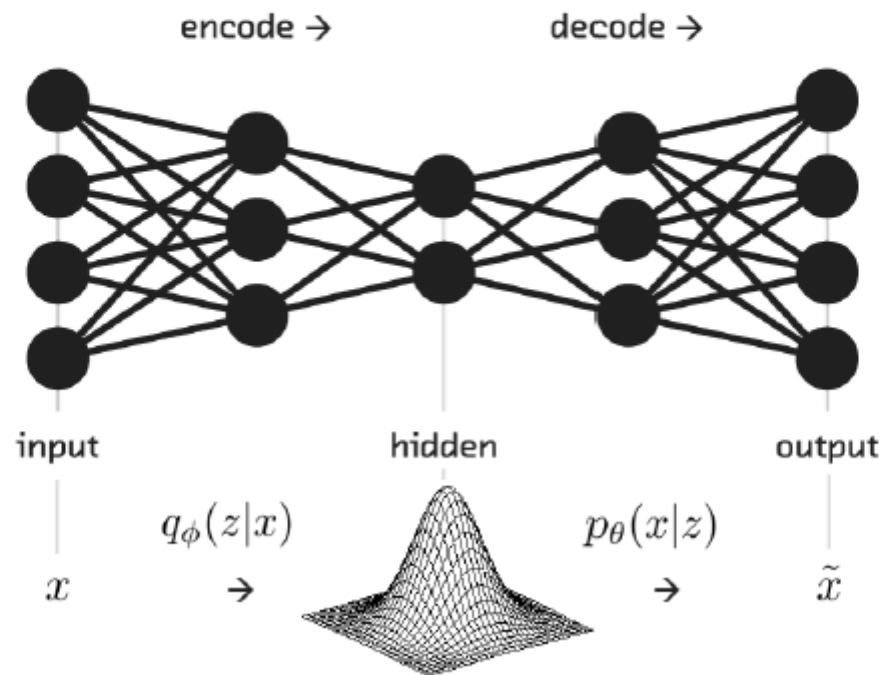
- Nonlinear hidden layer



- With non-linearity
 - “Non linear” PCA
 - Deeper networks can capture more complicated manifolds
 - “Deep” autoencoders

Variational Autoencoder (VAE)

- Objective $\mathcal{L}(x, \phi, \theta) = \underbrace{-D_{KL}(q_\phi(z|x)||p_\theta(z))}_{\text{Regularization term}} + \underbrace{E_{q_\phi(z|x)}[\log p_\theta(x|z)]}_{\text{Reconstruction term}}$

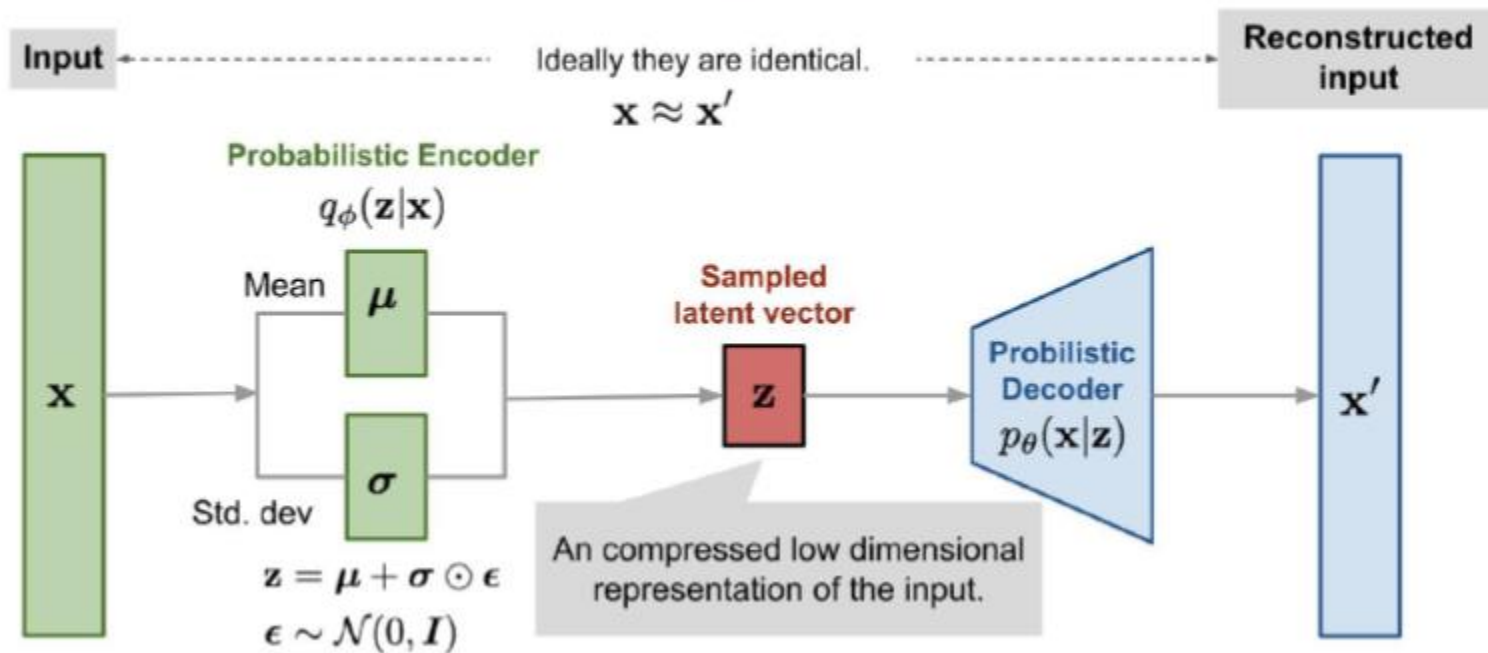


Variational Autoencoder (VAE)

■ Objective $\mathcal{L}(x, \phi, \theta) = -D_{KL}(q_{\phi}(z|x)||p_{\theta}(z)) + E_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)]$

Regularization term

Reconstruction term



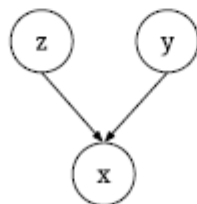
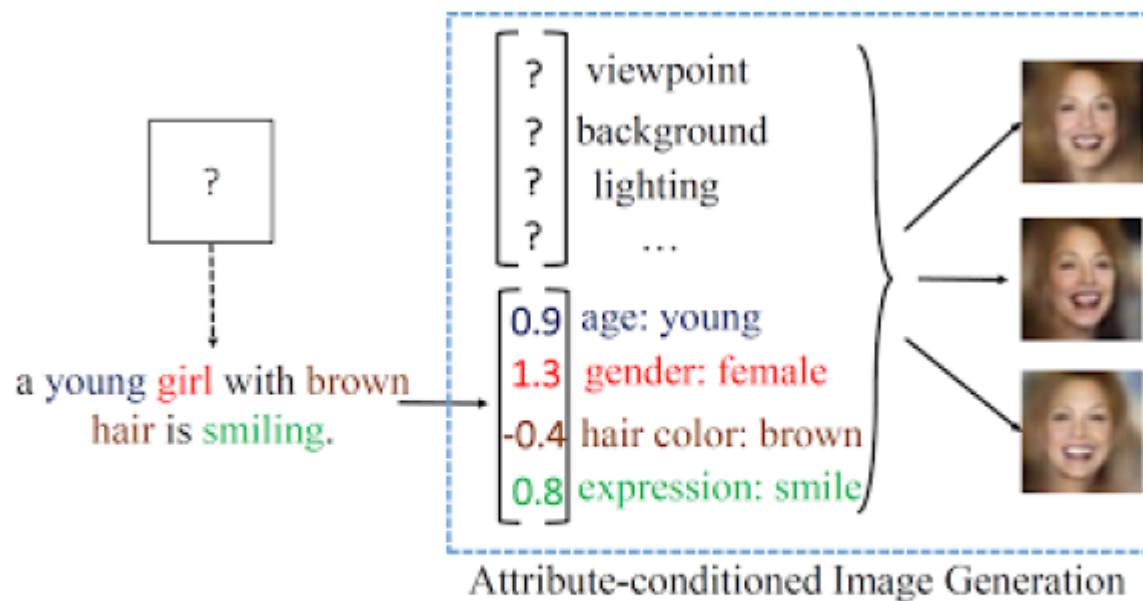
Interpreting the Latent Space



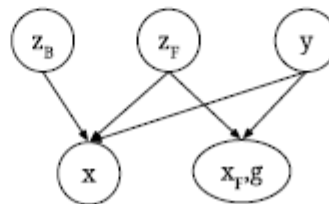
<https://arxiv.org/pdf/1610.00291.pdf>

Example: Attribute2Image

■ Main ideas



(a) CVAE: $p_{\theta}(x|y, z)$



(b) disCVAE: $p_{\theta}(x, x_F, g|y, z_F, z_B)$

Example: Attribute2Image

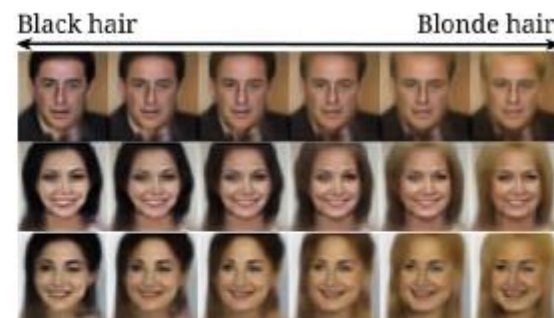
■ Results



(a) progression on gender



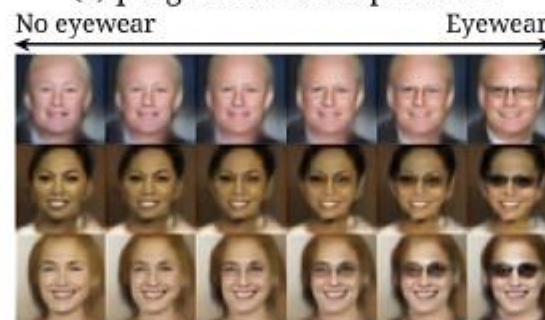
(c) progression on expression



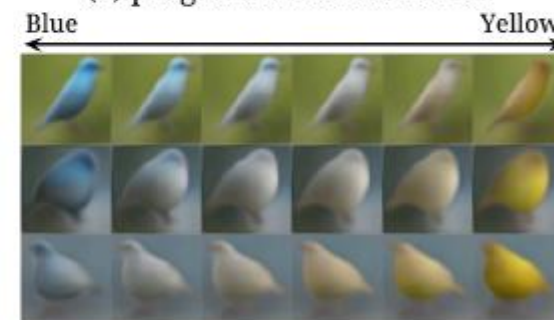
(e) progression on hair color



(b) progression on age



(d) progression on eyewear




(f) progression on primary color

Problem of VAE


- Blurry images



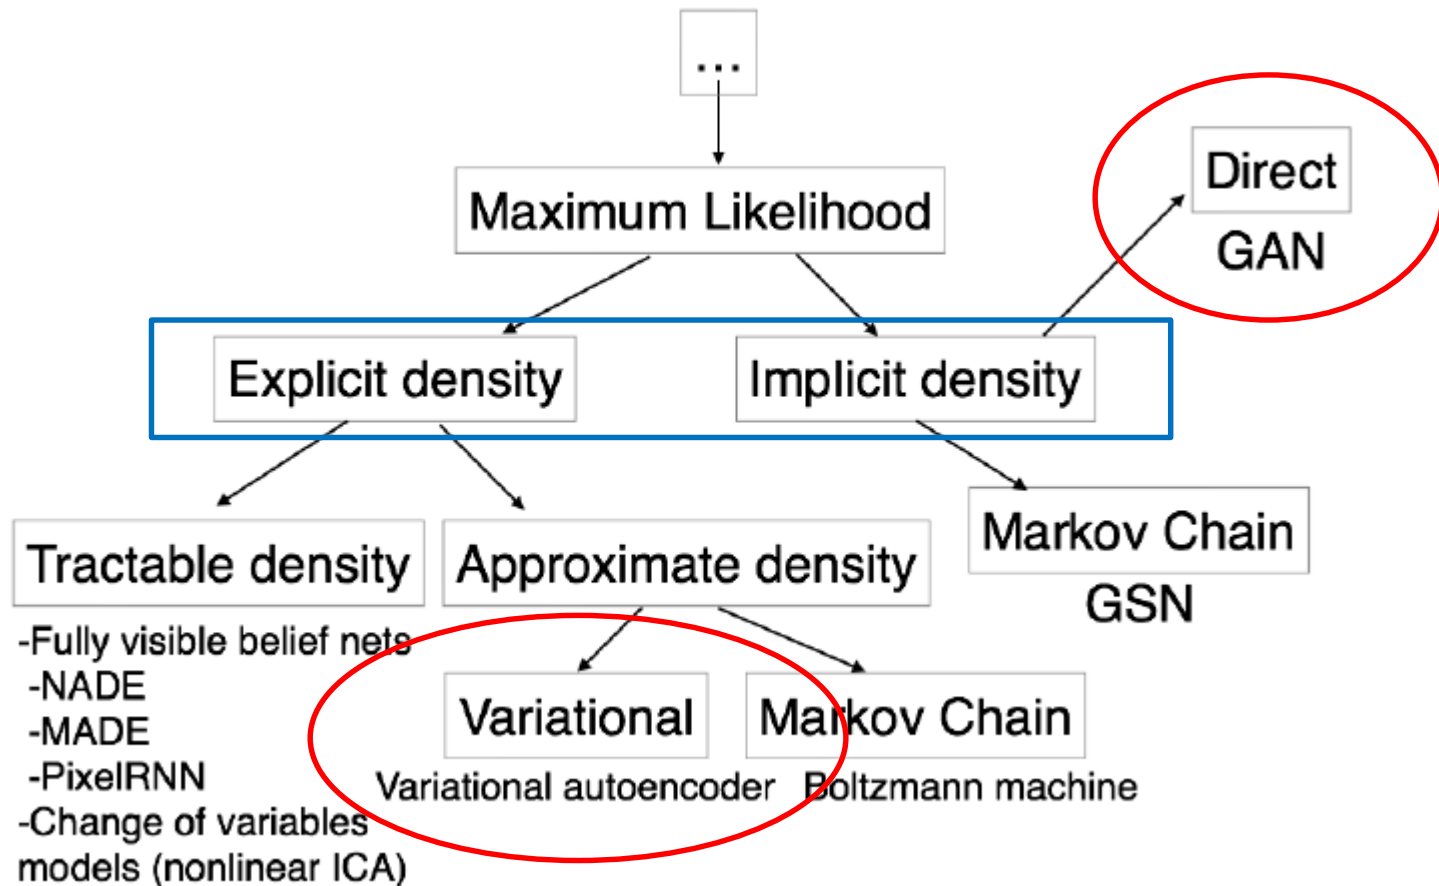
<https://blog.openai.com/generative-models/>



Deep Generative Networks-DGM

- Overview
 - Representation Learning with Autoencoder
 - **Generative Adversarial Network (GAN)**
 - Applications of GANs
- 

Taxonomy of Generative Models




Implicit Generative Models

- Working with explicit model $p(x)$ could be expensive
 - Variational Autoencoder (variational inference)
 - Boltzmann Machines (MCMC)
- Representation learning may not require $p(x)$
 - Sometimes we are more interested in taking samples from $p(x)$ instead of p itself

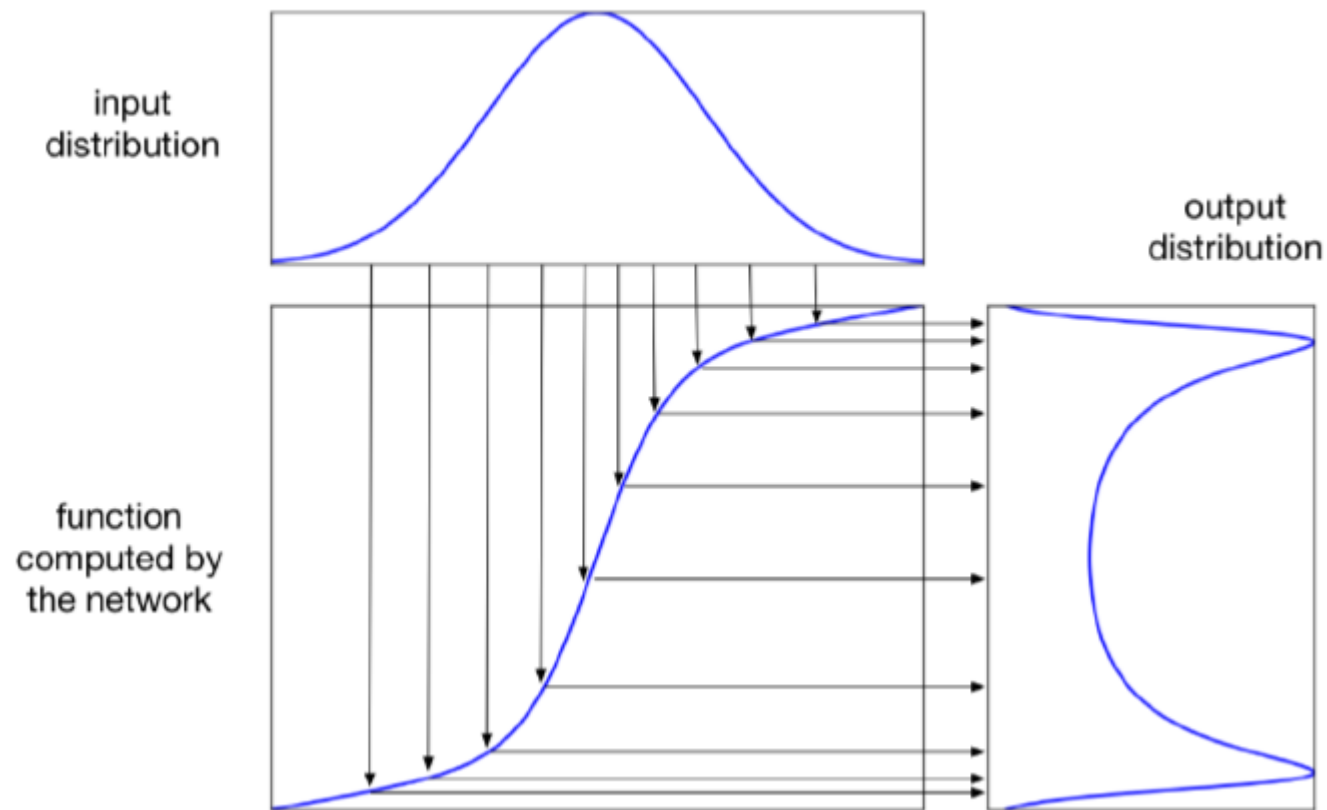
Implicit Generative Models

- Implicitly define a probability distribution
- Start by sampling the code vector z from a fixed, simple distribution
- A generator network computes a differentiable function G mapping z to an x in data space



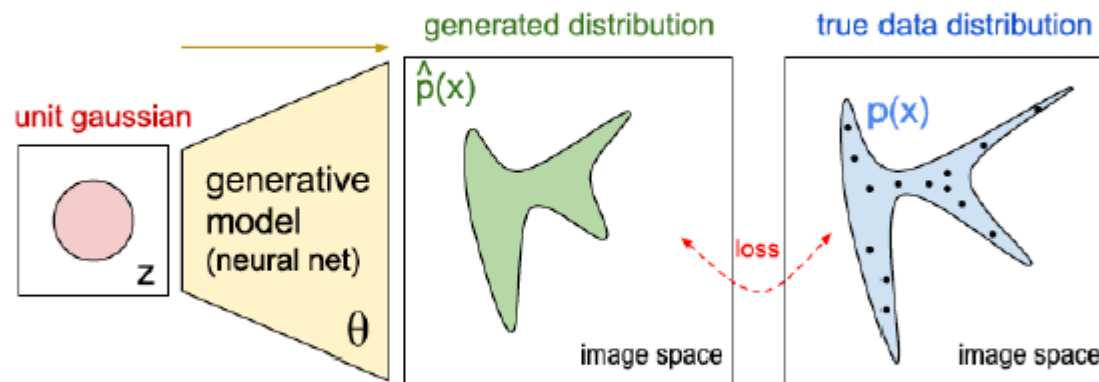
Implicit Generative Models

- Intuition: 1D example



Implicit Generative Models

■ Intuition

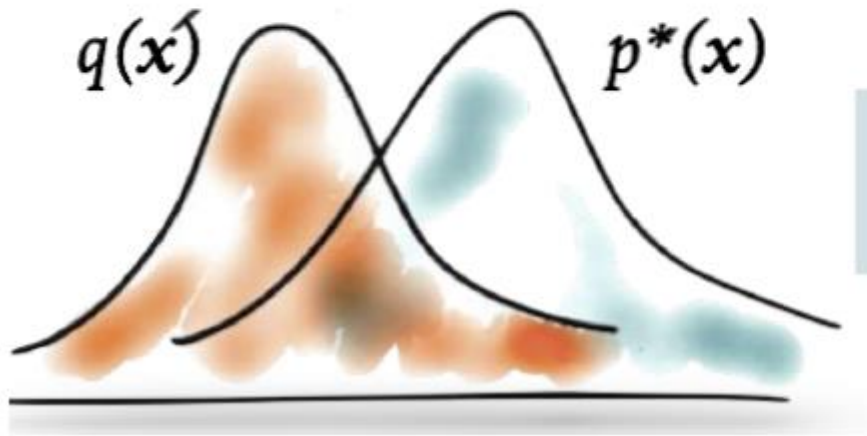


advocate/penalize samples within the blue/white region.

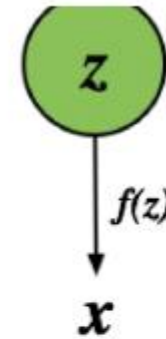
Learning by Comparison

■ Basic idea

For some models, we only have access to an unnormalised probability, partial knowledge of the distribution, or a simulator of data.



We compare the estimated distribution $q(x)$ to the true distribution $p^*(x)$ using samples.



Generative Adversarial Networks (GAN)

- Using a neural network to generate data

Output: Sample from
training distribution



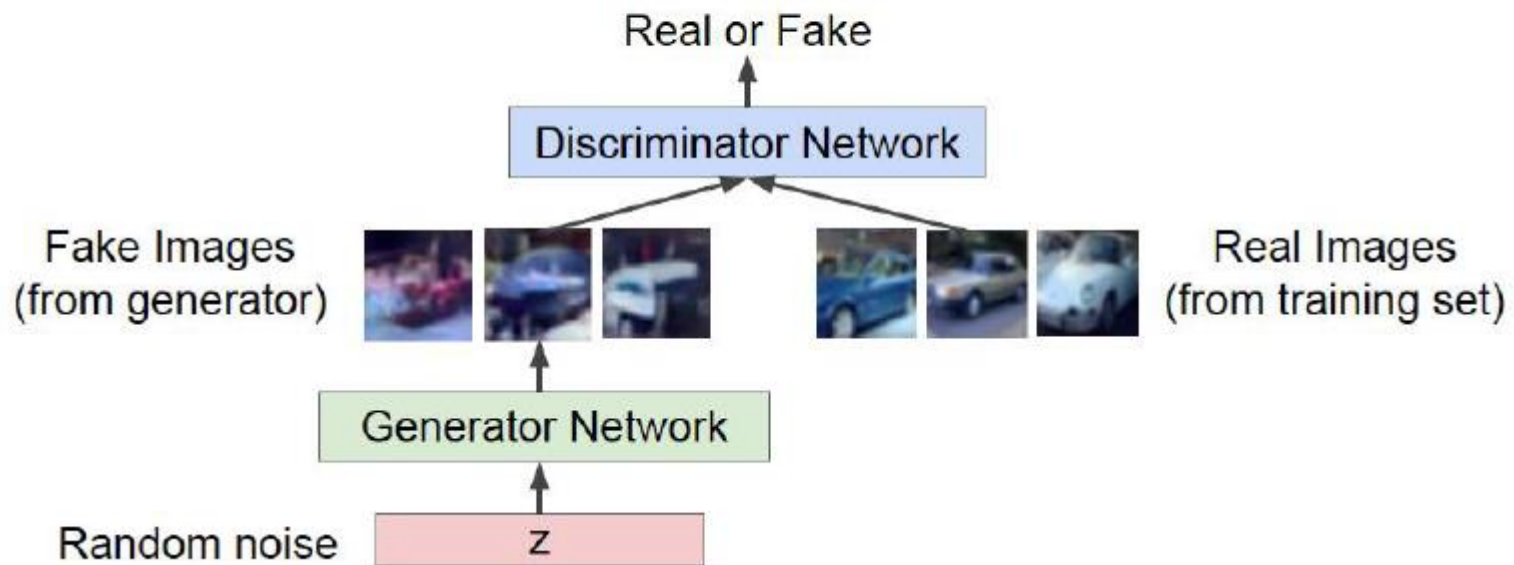
Generator
Network

Input: Random noise

z

Generative Adversarial Networks (GAN)

- Using another neural network to determine if the data is real or not




Adversarial Learning

- GAN objective for the generator is some complicated objective function defined by a neural network.
 - This means a new way of thinking about "distance".
 - We are training networks to minimize the "distance" or "divergence" between generated images and real images.
 - Instead of some hand-crafted distance metric like L1 or L2, we can make something completely new.
 - A neural network, with the right architecture, is arguably the definition of perceptual similarity (assuming our visual system is some sort of neural network).

Adversarial Learning

- Adversarial loss



Adversarial Learning

- Let D denote the discriminator's predicted probability of being real data
- Discriminator's cost function: cross-entropy loss for task of classifying real vs. fake images

$$\mathcal{J}_D = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[-\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z}}[-\log(1 - D(G(\mathbf{z})))]$$

- One possible cost function for the generator: the opposite of the discriminator's

$$\begin{aligned}\mathcal{J}_G &= -\mathcal{J}_D \\ &= \text{const} + \mathbb{E}_{\mathbf{z}}[\log(1 - D(G(\mathbf{z})))]\end{aligned}$$

Two-Player Game

■ Minimax formulation

- The generator and discriminator are playing a zero-sum game against each other

$$\max_G \min_D \mathcal{J}_D$$

- Using parametric models

Discriminator outputs likelihood in (0,1) of real image

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

Learning Procedure

■ Minimax objective function

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

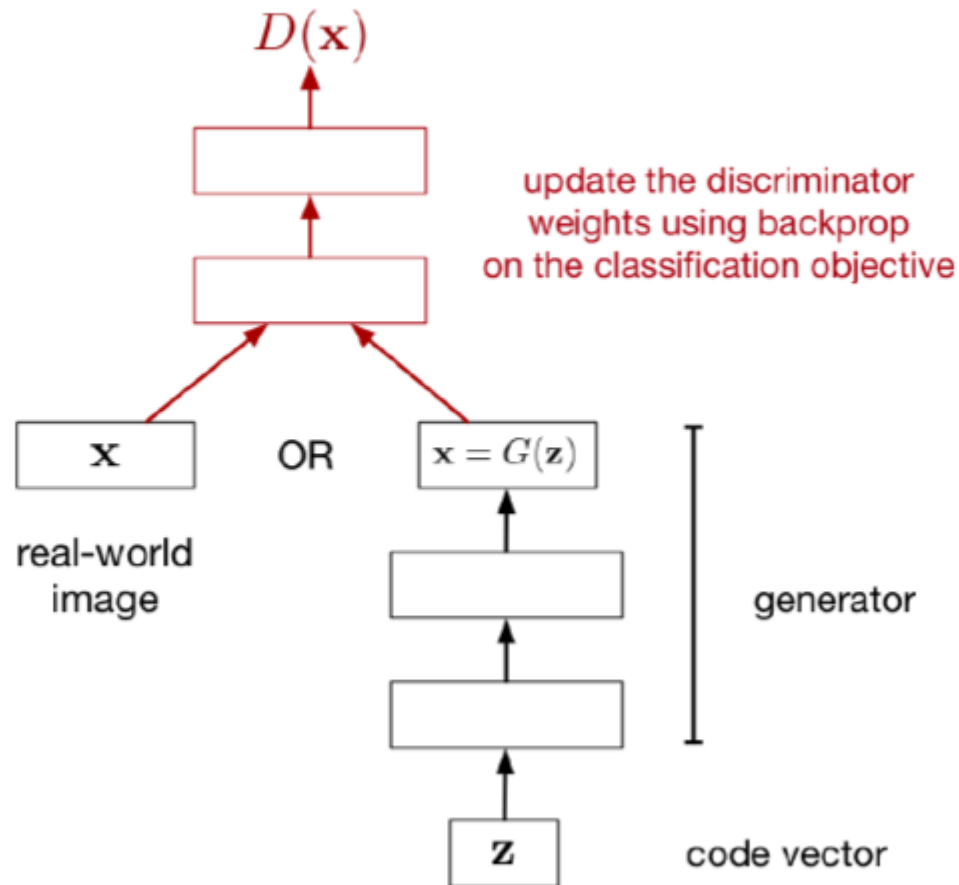
$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Gradient descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

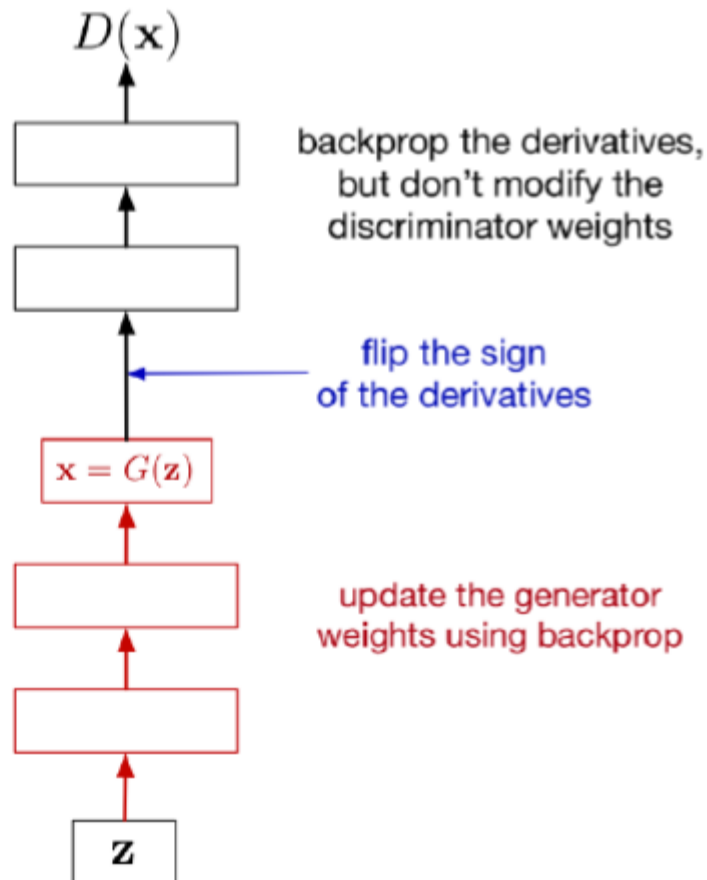
Learning Procedure

- Updating the discriminator

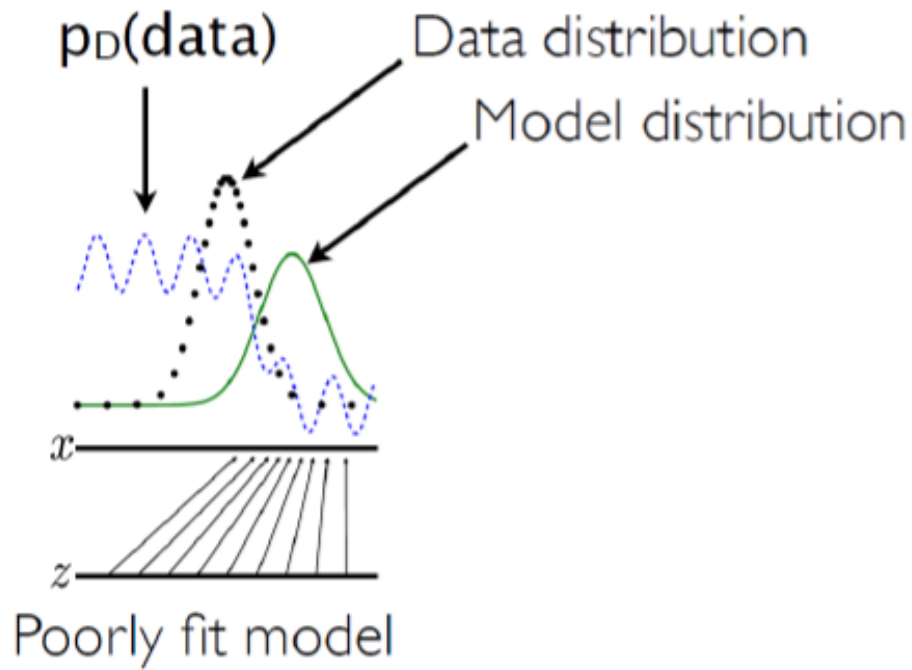


Learning Procedure

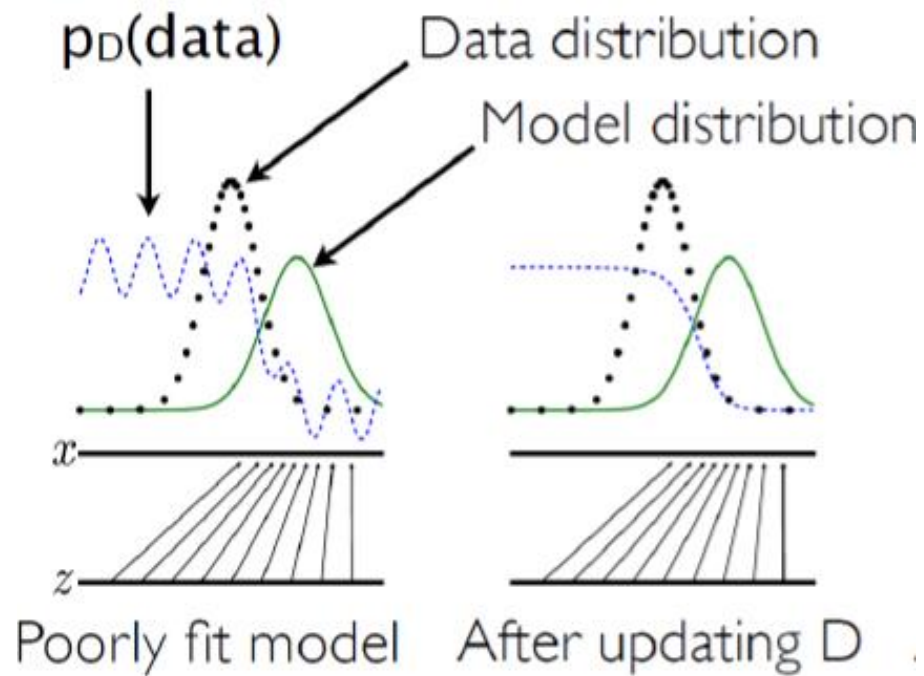
■ Updating the generator



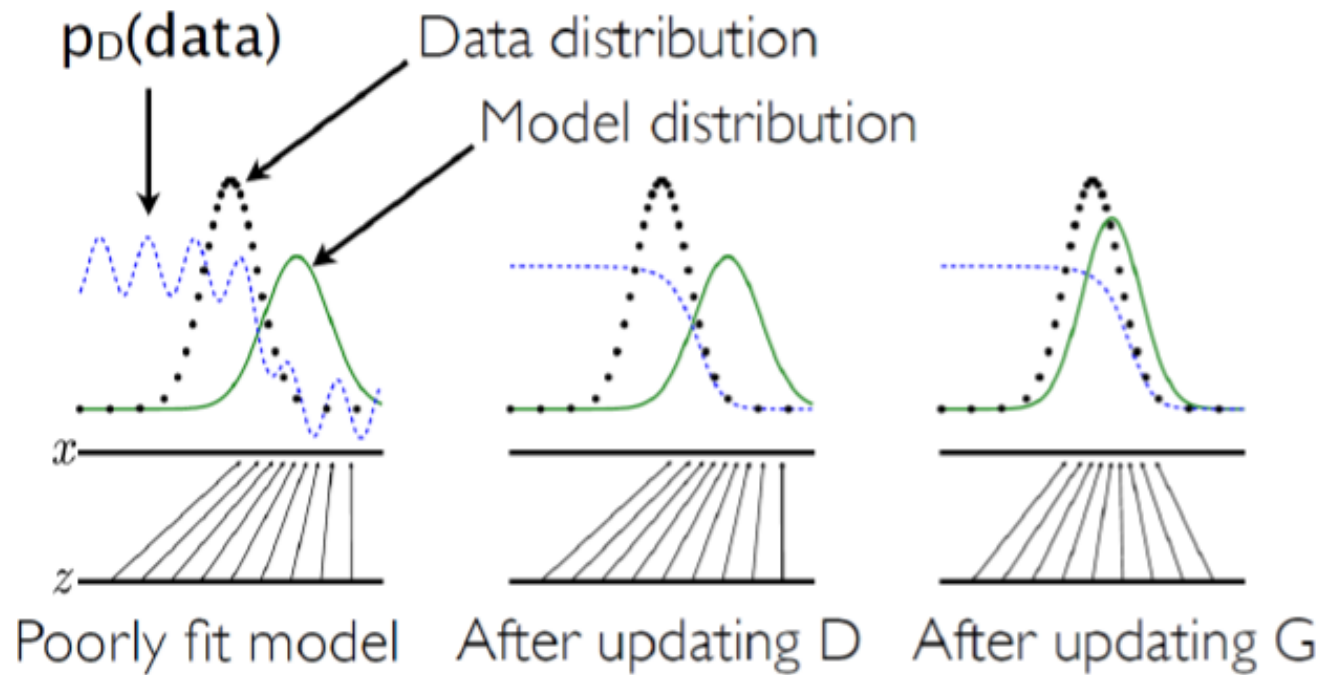
Training GANs



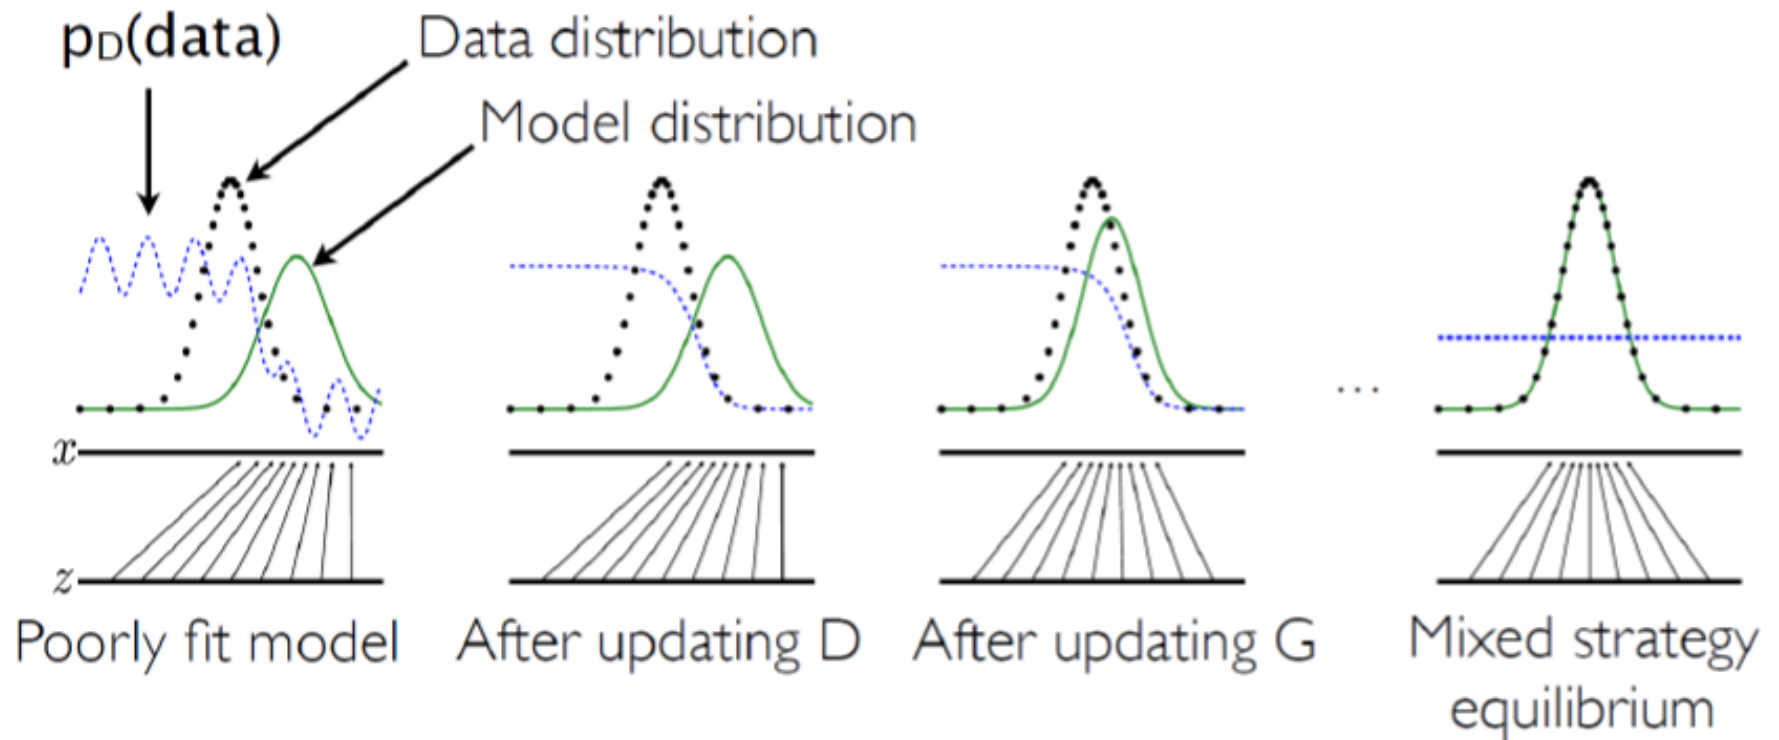
Training GANs



Training GANs

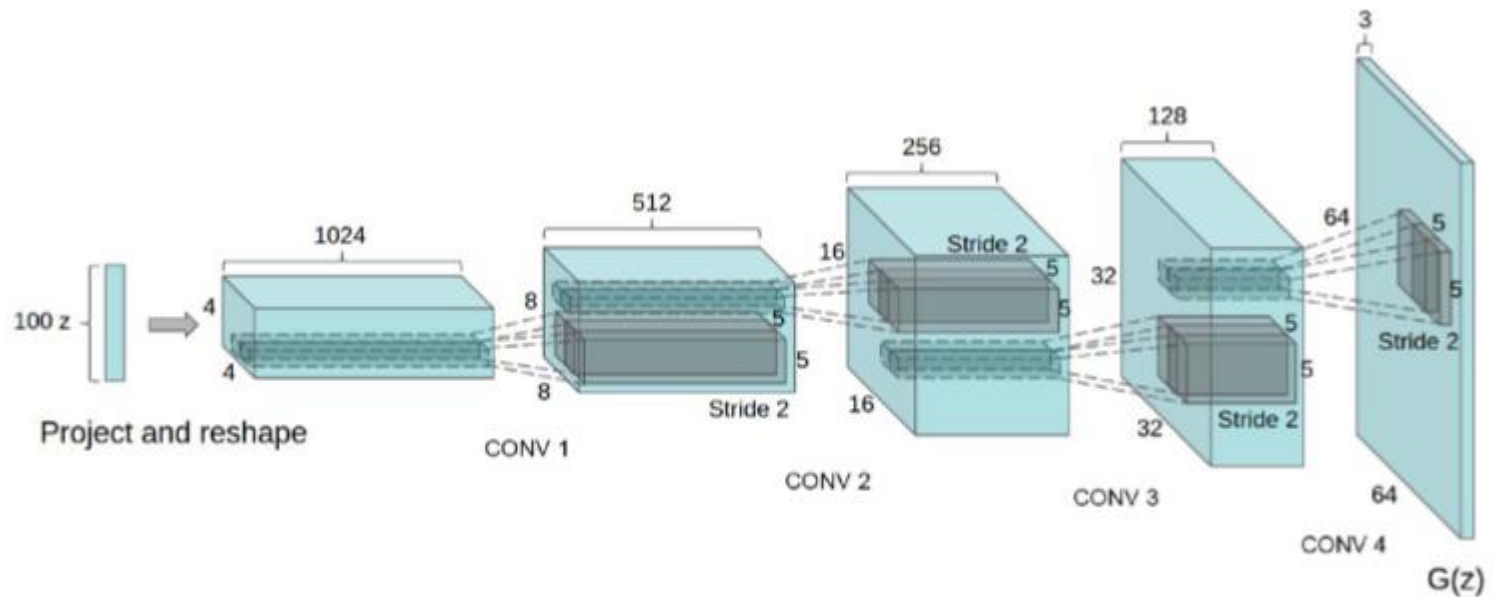


Training GANs



Typical Generator Architecture

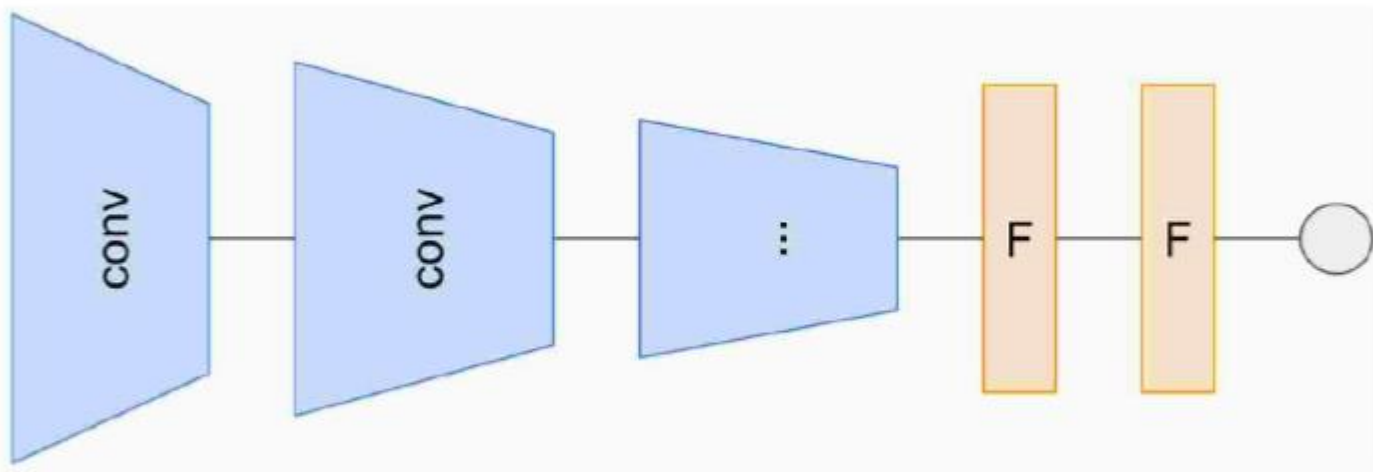
■ For images



- ▶ Unit Gaussian distribution on z , typically 10-100 dim.
- ▶ Up-convolutional deep network (reverse recognition CNN)

Typical Discriminator Architecture

- For images



- ▶ Recognition CNN model
- ▶ Binary classification output: real / synthetic

Training GANs

- Since GANs were introduced in 2014, there have been hundreds of papers introducing various architectures and training methods
- GAN Zoo: <https://github.com/hindupuravinash/the-gan-zoo>
- In general, training a GAN is tricky and unstable
- Many tricks:
 - S. Chintala, How to train a GAN, ICCV 2017 tutorial
 - https://github.com/soumith/talks/blob/master/2017-ICCV_Venice/How_To_Train_a_GAN.pdf

Generative Samples

Celebrities:



Karras et al., 2017. Progressive growing of GANs for improved quality, stability, and variation

Generative Samples

Objects:



Walk Around Data Manifold

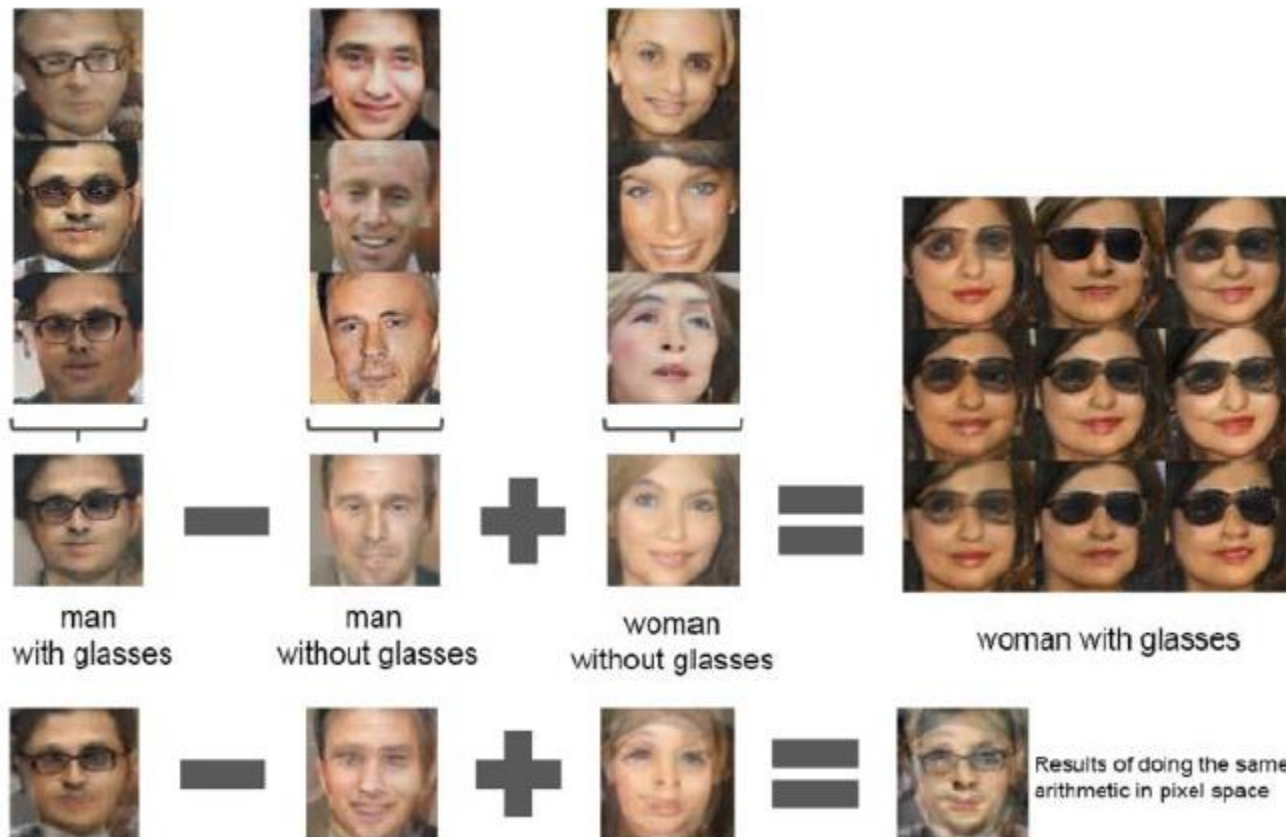
Interpolating
between
random
points in laten
space


Radford et al,
ICLR 2016




Walk Around Data Manifold

■ Vector Arithmetic





Deep Generative Networks-DGM

- Overview
 - Representation Learning with Autoencoder
 - Generative Adversarial Network (GAN)
 - **Applications of GANs**
- 

Conditional GANs

- Conditional GANs include a label and learn $P(X|Y)$
 - Add conditional variable y into G and D
 - Objective function

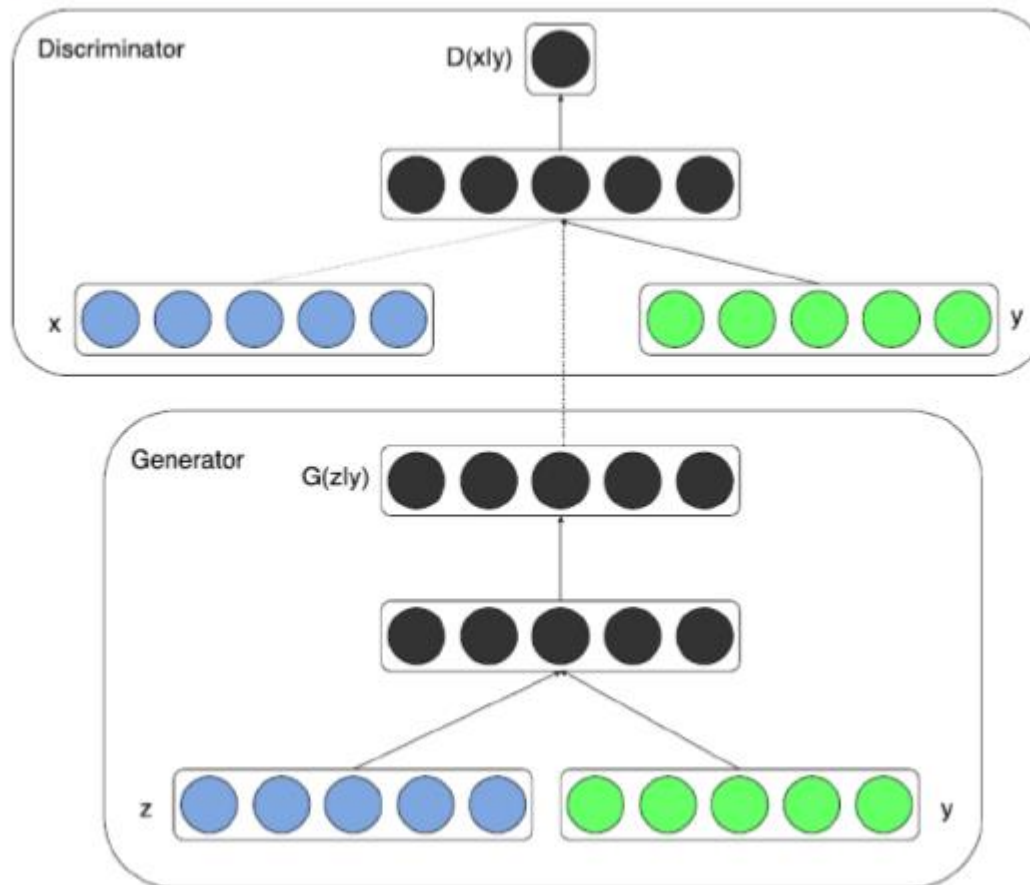
$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$



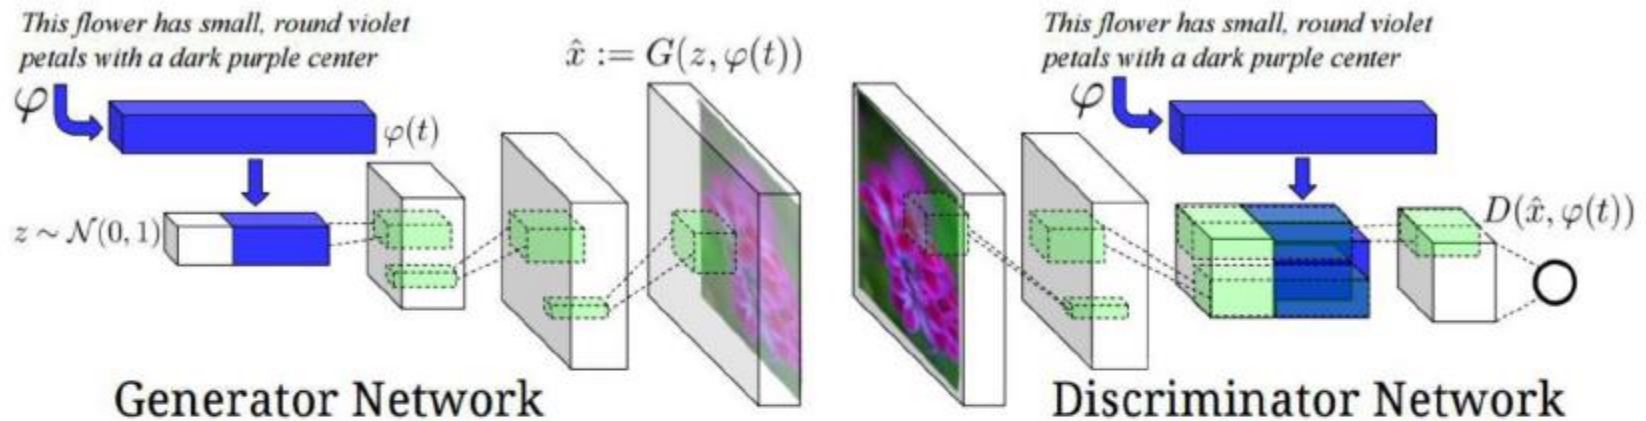
$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|\mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|\mathbf{y})))].$$

Conditional GANs

■ Model architecture



Conditional GANs



this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



Reed et al 2015

StackGAN

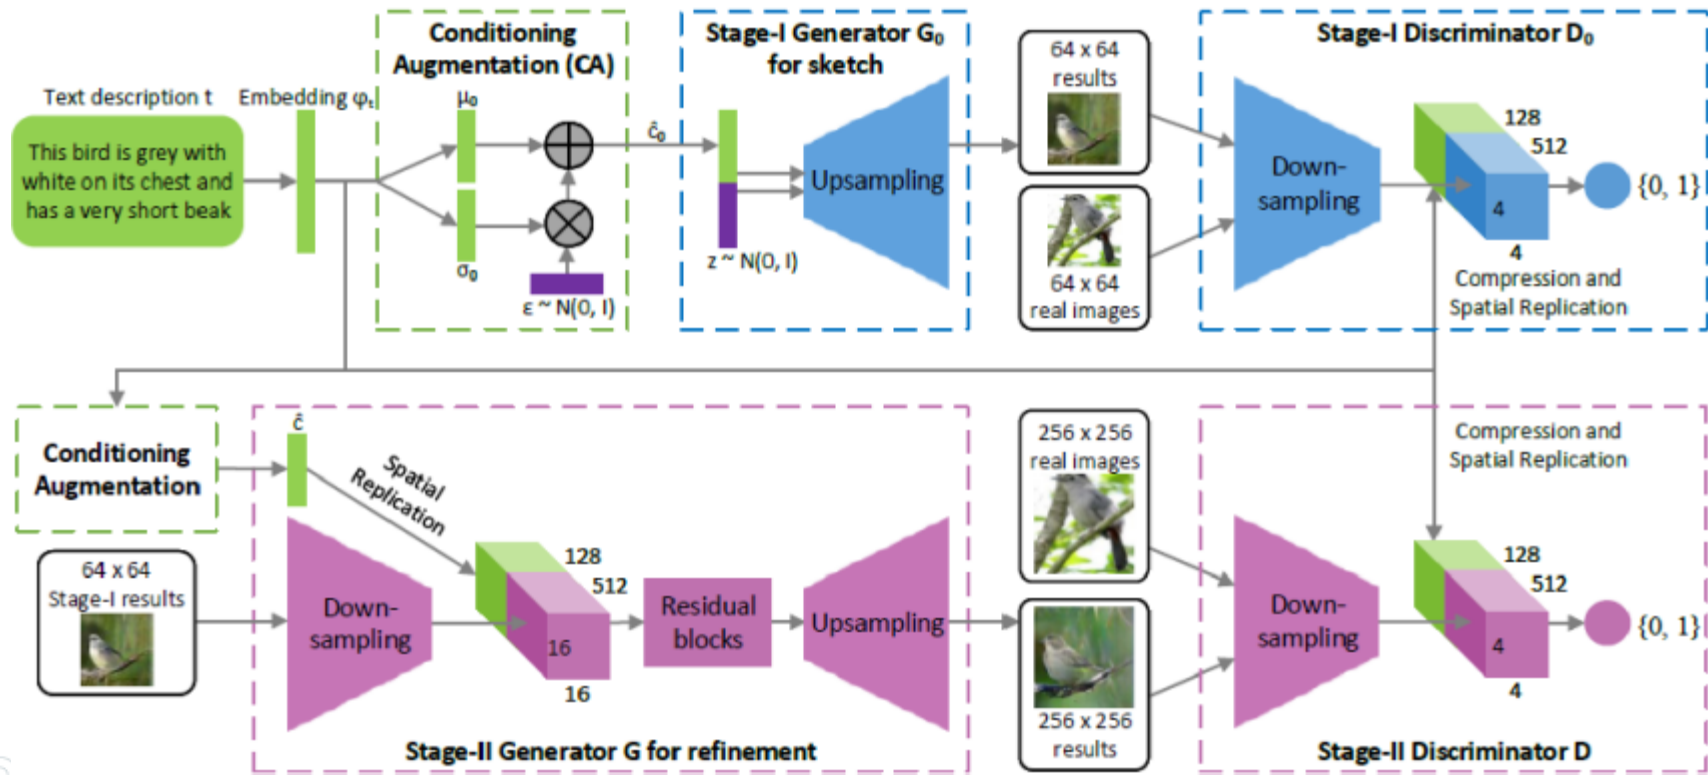
■ A coarse-to-fine manner



Zhang et al. 2016

StackGAN

- Use stacked GAN structure



More StackGAN results

Text description	This flower is pink, white, and yellow in color, and has petals that are striped	This flower has a lot of small purple petals in a dome-like configuration	This flower is white and yellow in color, with petals that are wavy and smooth	This flower has petals that are dark pink with white edges and pink stamen
64x64 GAN-INT-CLS				
256x256 StackGAN				

Image-to-Image Translation

- One-to-many or many-to-one mapping [Isola et al., 2016]

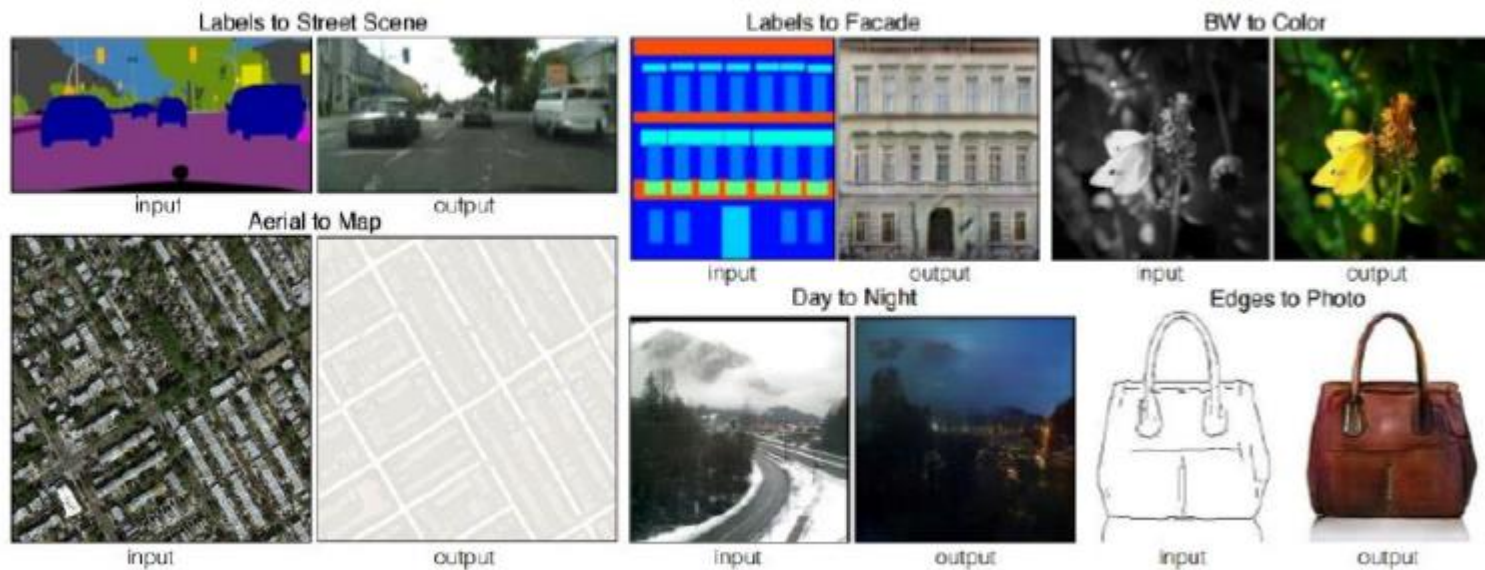


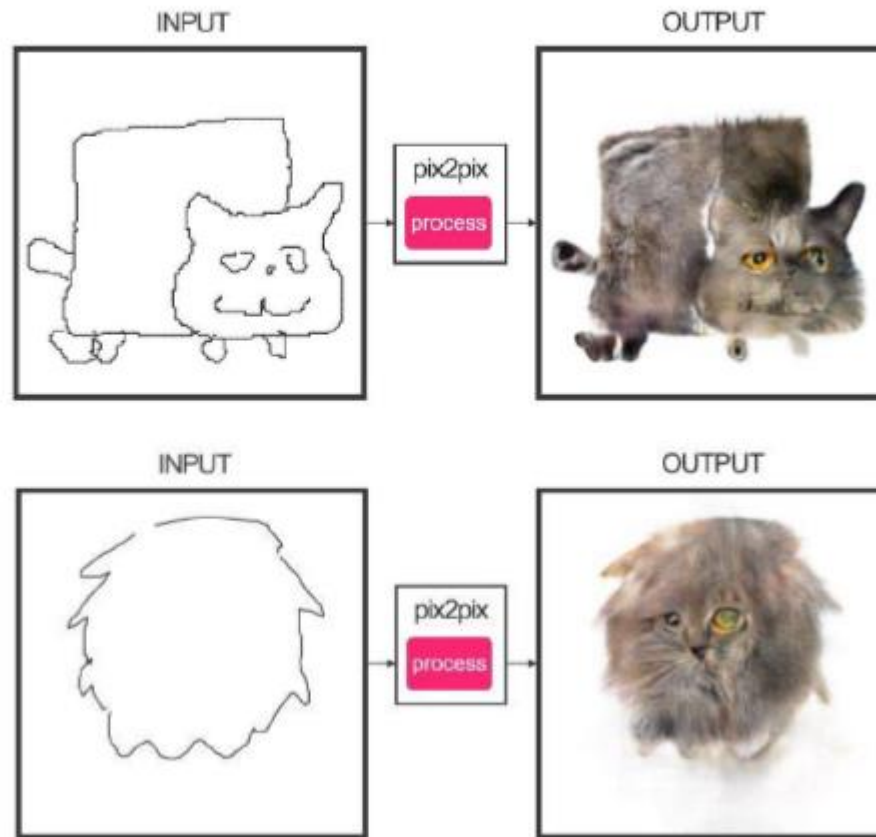
Image-to-Image Translation

■ More results



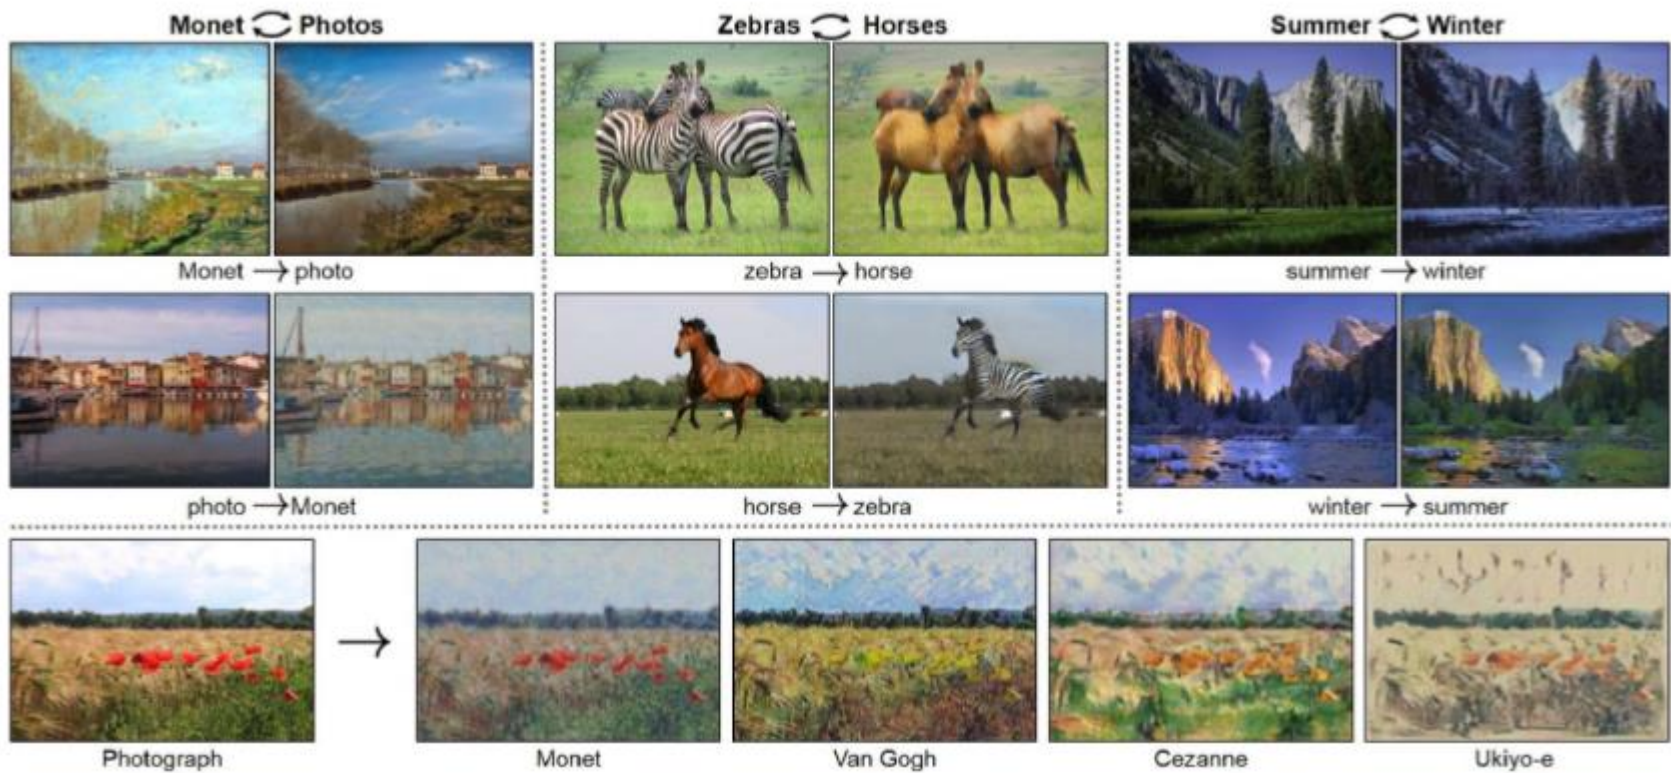
Image-to-Image Translation

■ More results



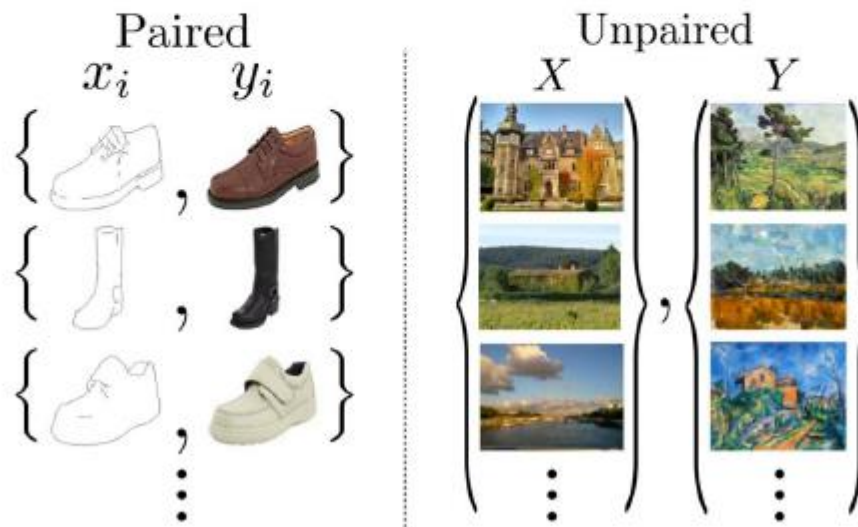
CycleGAN

- Image-to-image translation without paired data



CycleGAN

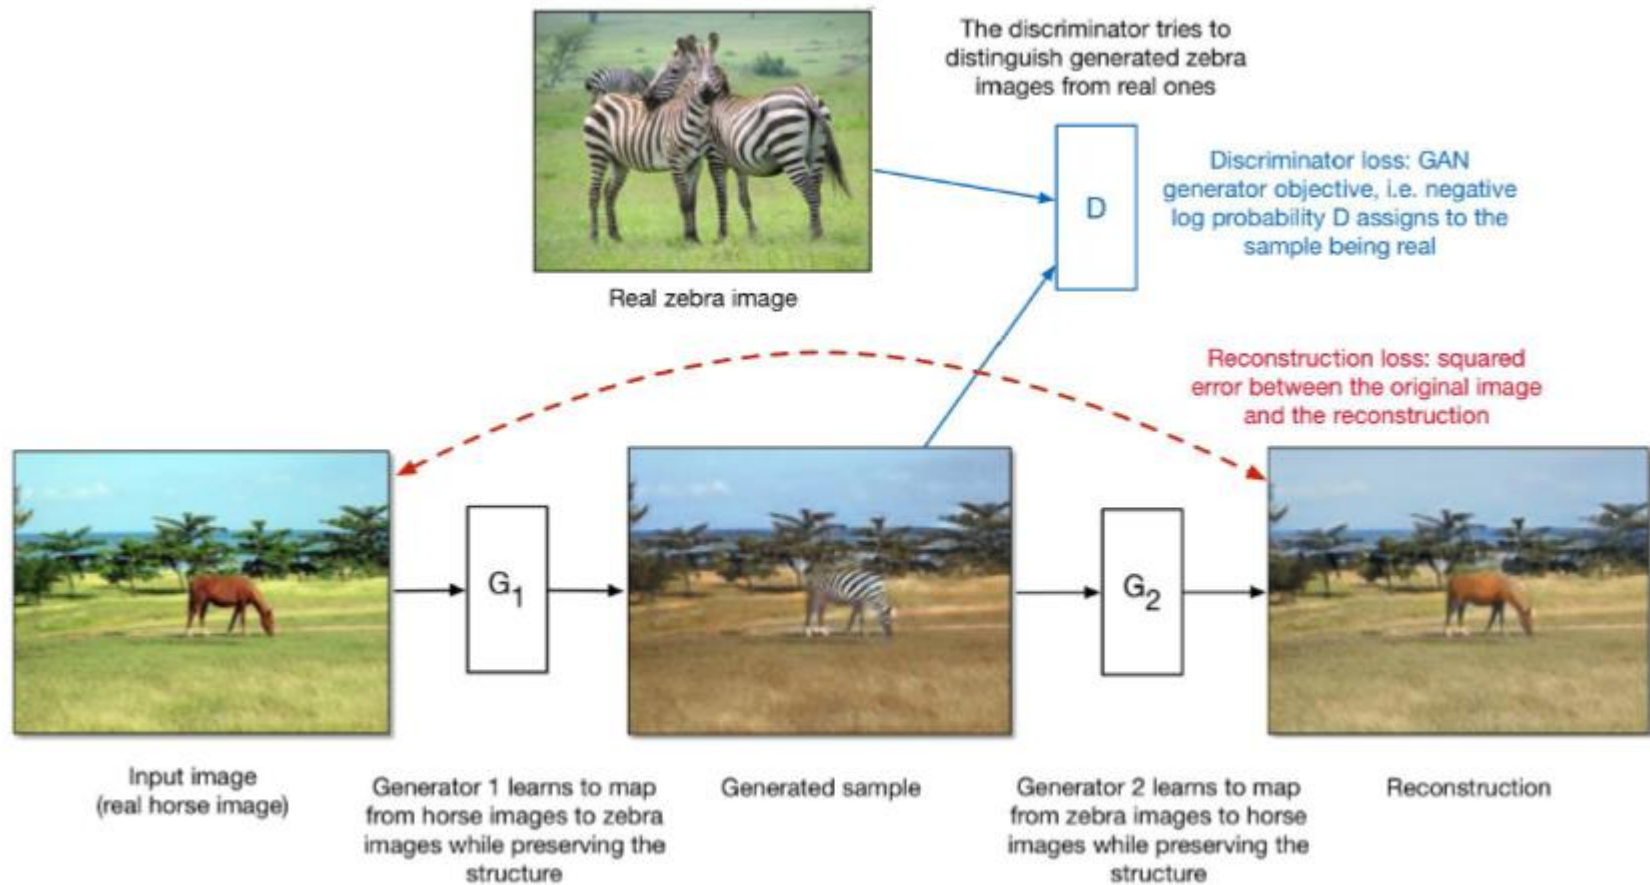
- If we had paired data (same content in both styles), this would be a supervised learning problem. But this is hard to find.



CycleGAN

- If we had paired data (same content in both styles), this would be a supervised learning problem. But this is hard to find.
- The CycleGAN architecture learns to do it from unpaired data.
 - Train two different generator nets to go from style 1 to style 2, and vice versa.
 - Make sure the generated samples of style 2 are indistinguishable from real images by a discriminator net.
 - Make sure the generators are **cycle-consistent**: mapping from style 1 to style 2 and back again should give you almost the original image.

CycleGAN



$$\text{Total loss} = \text{discriminator loss} + \text{reconstruction loss}$$

CycleGAN

■ Results



■ More details

<https://hardikbansal.github.io/CycleGANBlog/>